



Module 4 : Les Architectures Conteneurisées

16 & 17 mars 2023

Thomas Lecler
Software Engineer
Devops



Amaury Bloch
Software Engineer
Devops



QUI SUIS-JE ?

Thomas
LECLER

Software Engineer - DevOps

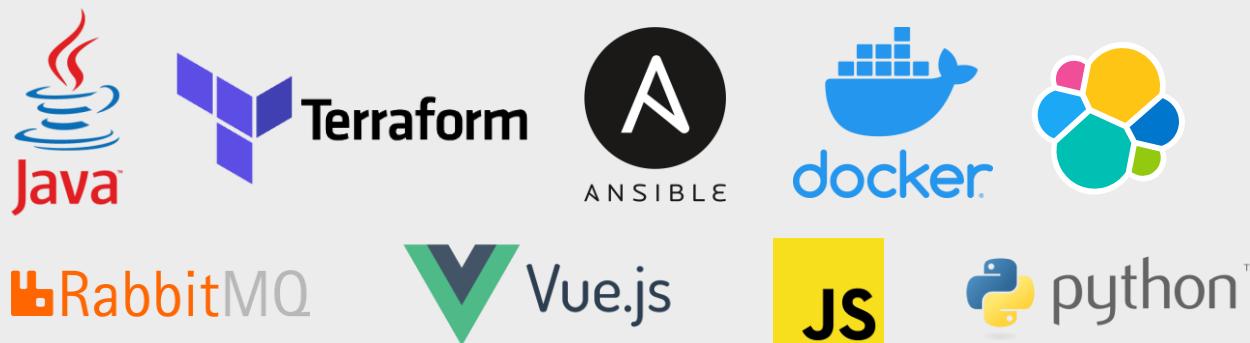
Mon expertise

4 Projets (Ministère des Armées / Ministère de l'Intérieur / Ministère de l'Agriculture)

Mon parcours

- A Capgemini depuis le Stage de Fin d'Étude
- Avant Capgemini :
 - École d'ingénieur en Informatique, Réseau et Cybersécurité
 - CPGE : Mathématiques et Physique

Mes technos



QUI SUIS-JE ?

Amaury
BLOCH

Software Engineer - DevOps

Mon expertise

6 Projets (Ministère des Armées / Ministère de l'Agriculture)

Mon parcours

- A Capgemini depuis le Stage de Fin d'Étude
- Avant Capgemini :
 - École d'ingénieur généraliste et spécialité informatique
 - CPGE : Physique, Technologie et Sciences de l'ingénieur

Mes technos



kubernetes



docker



ANSIBLE



CI⁺CD



RabbitMQ



Vue.js



spring[®]



AGENDA

- Module 1 – Le Cloud Computing
Comprendre de manière critique le Cloud
- Module 2 – Les Architectures Cloud
Instancier une architecture reposant sur les services présentés
- Module 3 – Scalabilité
- Module 4 – Les Architectures Conteneurisées
Réduire l'adhérence vis-à-vis des clouders

- Module 4 – Les Architectures Conteneurisées
Réduire l'adhérence vis-à-vis des clouders

Conteneuriser ses applications

- Concepts & bénéfices
- Création d'une image
- Lancement d'un conteneur

Orchestrer ses conteneurs via Kubernetes

- Concepts & bénéfices
- Déployer un cluster Kubernetes
- Les ressources Kubernetes
- Déployer des ressources sur le cluster

Industrialiser ses déploiements

- Du code au cluster
- Automatiser la construction et le déploiement

Exploiter & superviser ses applications en production

- Exploiter ses applications
- Superviser ses applications
- Collecter des métriques

A photograph of a shipping container yard at sunset. In the foreground, a red and black straddle carrier crane is positioned between stacks of shipping containers. One blue shipping container is being lifted by the crane's hook. The background shows a massive stack of shipping containers in various colors, including blue, red, and green, stacked in a grid pattern against a sky filled with warm orange and yellow hues.

CONTENEURISER SES APPLICATIONS POUR LES RENDRE PORTABLES



Conteneuriser ses applications

- Concepts & bénéfices
- Création d'une image
- Lancement d'un conteneur

Orchestrer ses conteneurs via Kubernetes

- Concepts & bénéfices
- Déployer un cluster Kubernetes
- Les ressources Kubernetes
- Déployer des ressources sur le cluster

Industrialiser ses déploiements

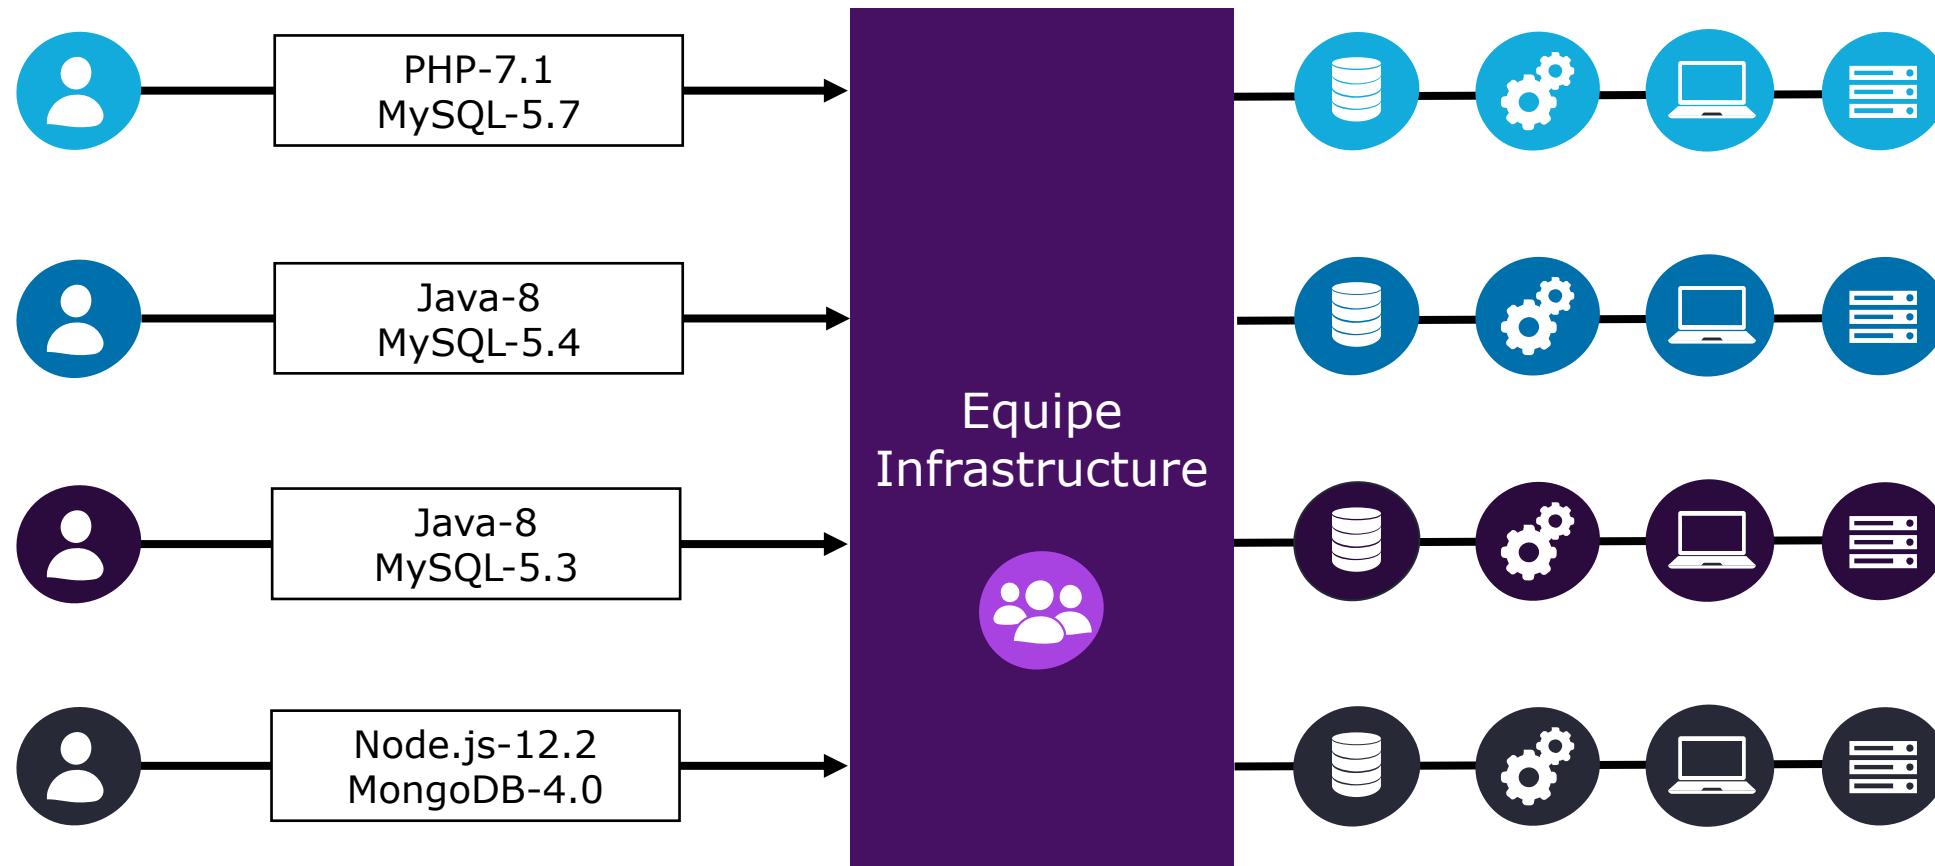
- Du code au cluster
- Automatiser la construction et le déploiement

Exploiter & superviser ses applications en production

- Exploiter ses applications
- Superviser ses applications
- Collecter des métriques

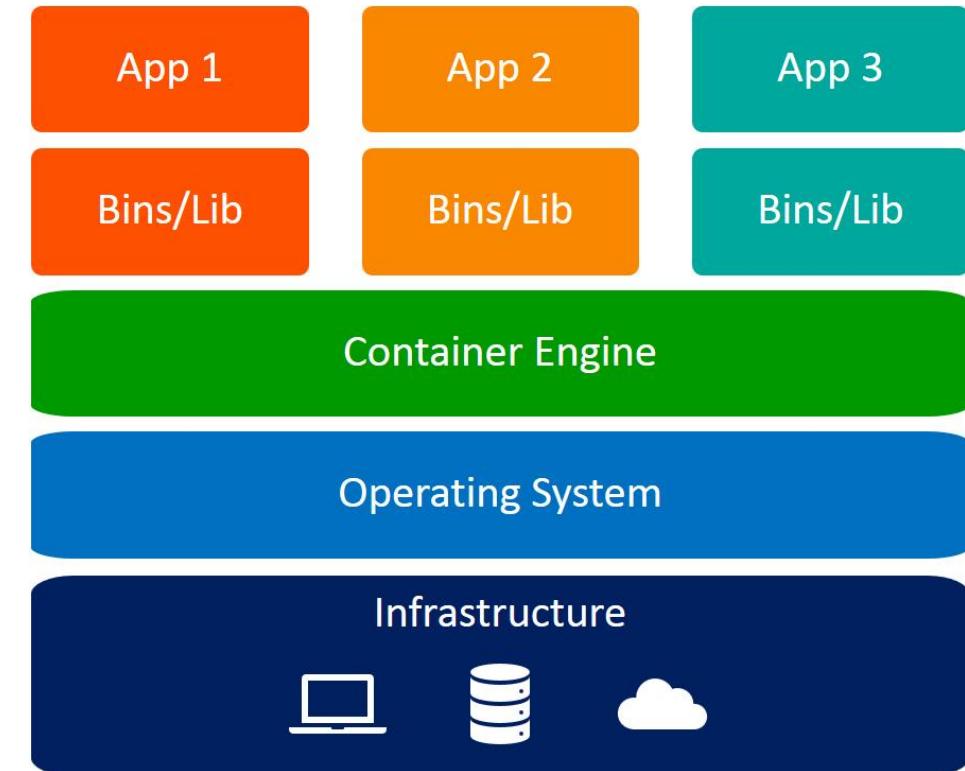
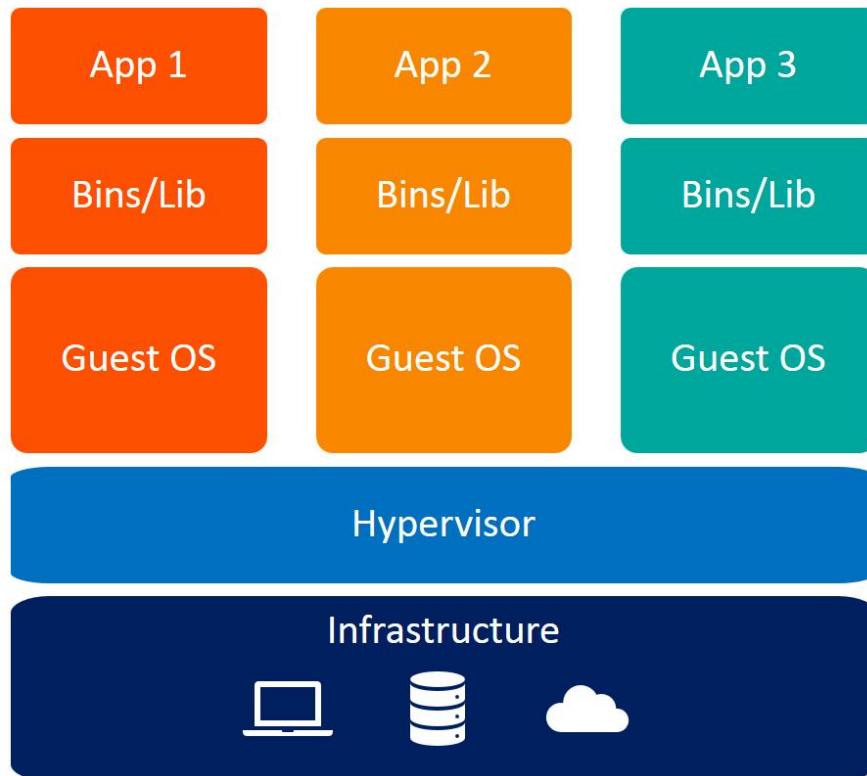
CONCEPTS

Déployer une application en entreprise



CONCEPTS

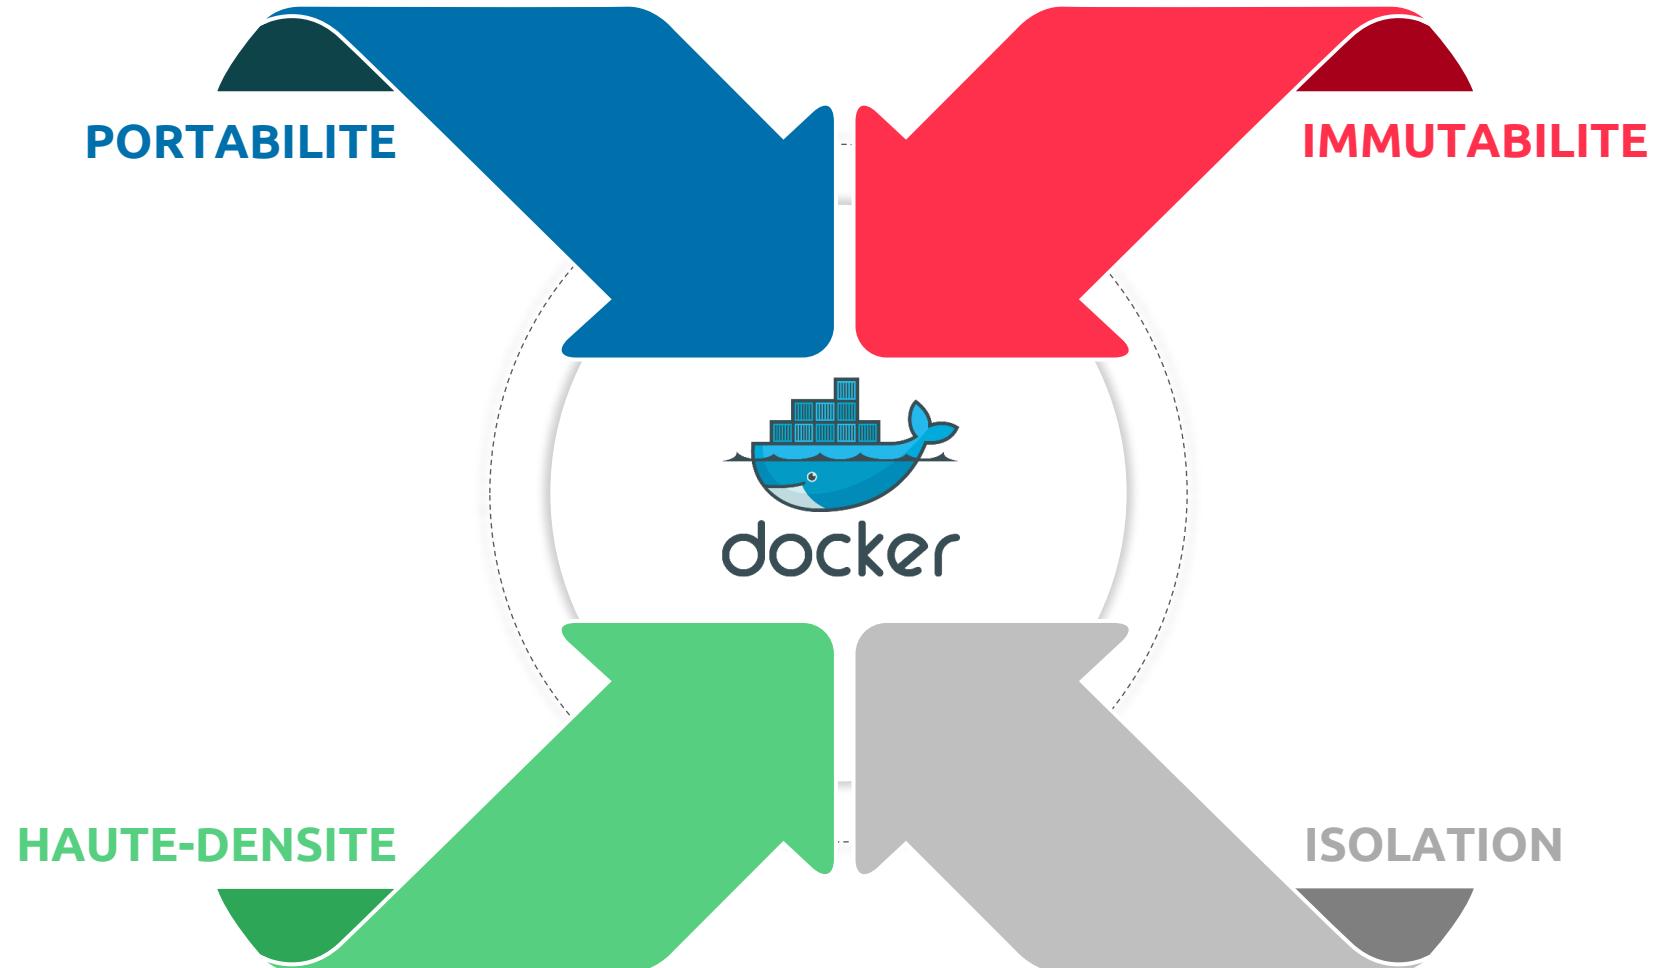
Virtualisation vs Conteneurisation



<https://www.bmc.com/blogs/containers-vs-virtual-machines>

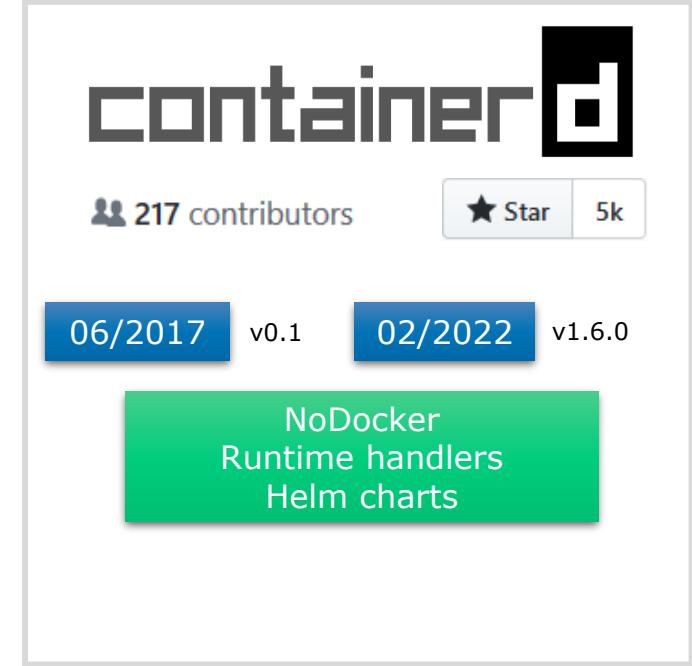
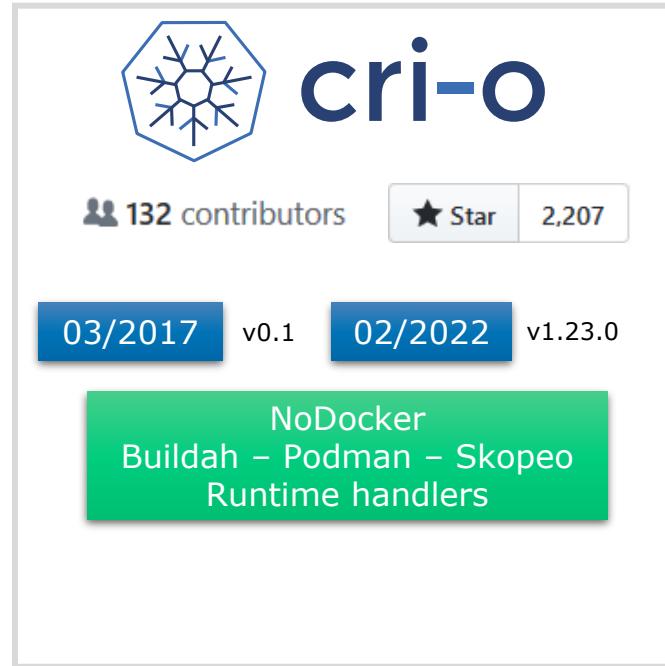
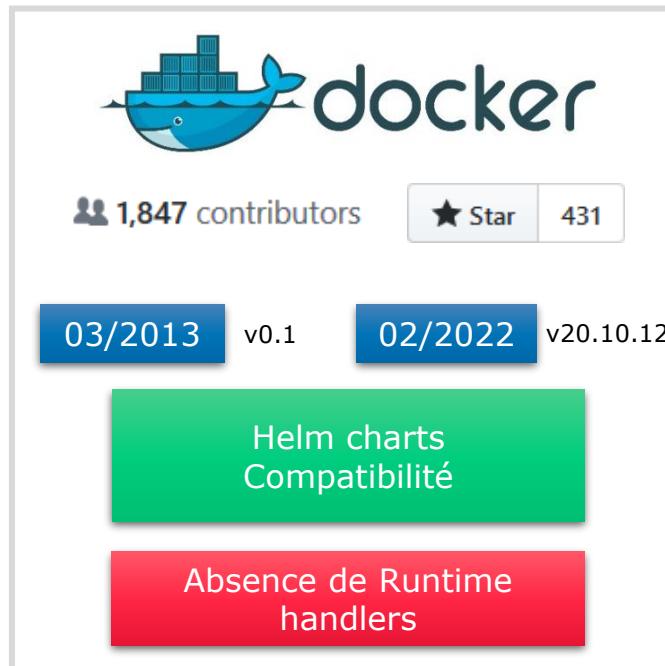
CONCEPTS

Bénéfices de la conteneurisation



CONTENEURISATION

Choix du Container Engine



Opportunités



Sandboxing



Multi Container Runtime

LES STANDARDS DE LA CONTENEURISATION

Open Container Initiative et Container Runtime Initiative

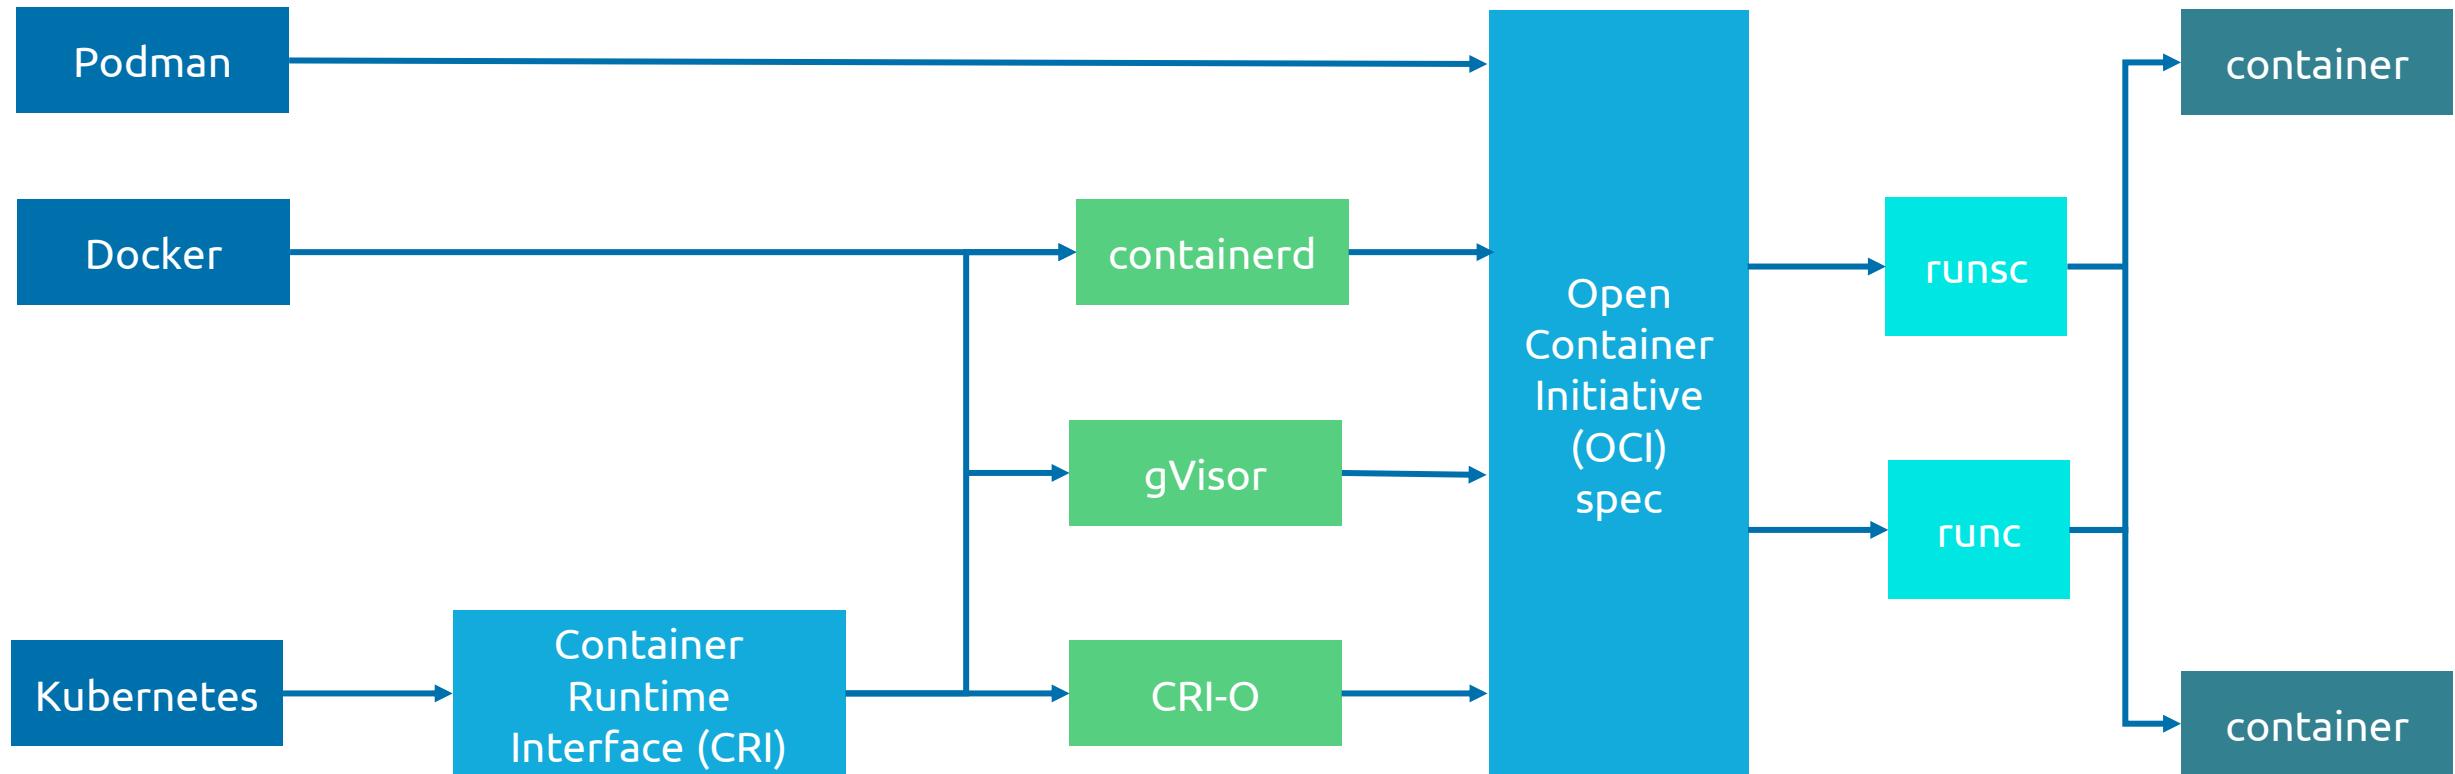
L'**Open Container Initiative (OCI)** publie des spécifications pour les conteneurs et leurs images.

La Kubernetes **Container Runtime Interface (CRI)** définit une API entre Kubernetes et un conteneur runtime sous-jacent.

OCI permet de choisir la chaîne d'outils pour conteneurs la plus adaptée au contexte (e.g. Docker, CRI-O, Podman, LXC, etc).

LES ACTEURS DU DÉPLOIEMENT D'UN CONTENEUR

Construire et exécuter des conteneurs

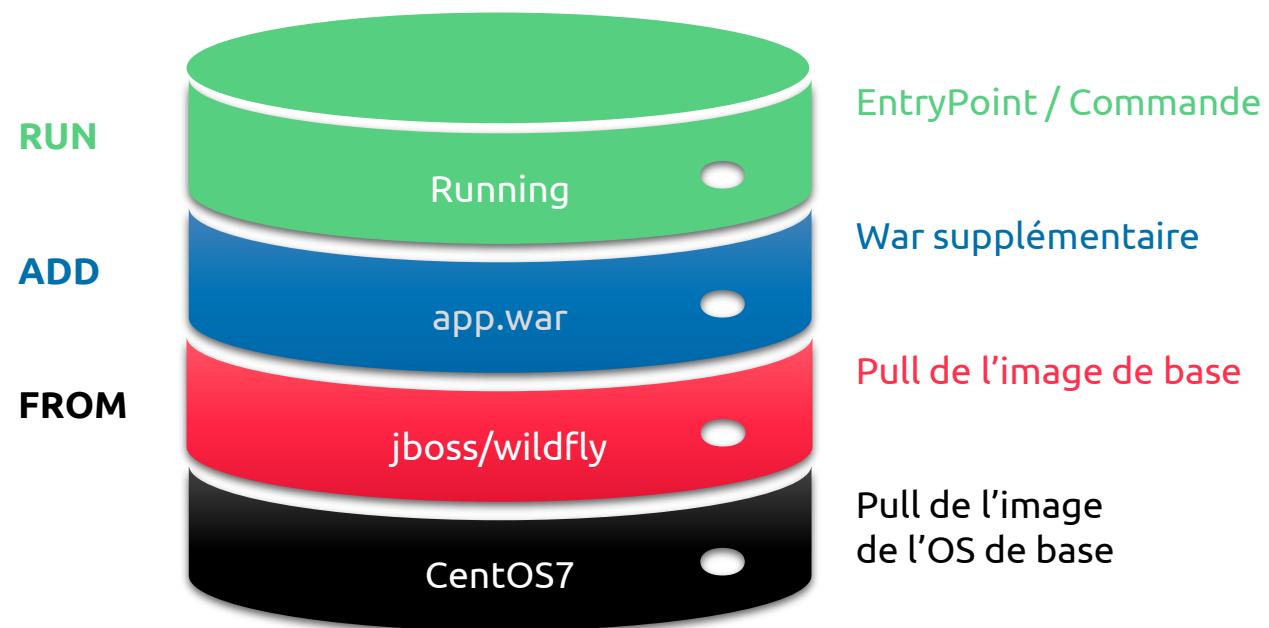


CRÉATION D'UNE IMAGE

L'image : version statique de l'application

Dockerfile

```
FROM jboss/wildfly:11  
  
ADD /target/*.war  
/opt/jboss/wildfly/standalone/deploy  
ments/  
  
RUN echo 'we are running some # of  
cool things'  
  
EXPOSE 8080  
USER 9000
```



HÉBERGER SON APPLICATION

Utiliser des registres publics ou on-premise

Registres publics

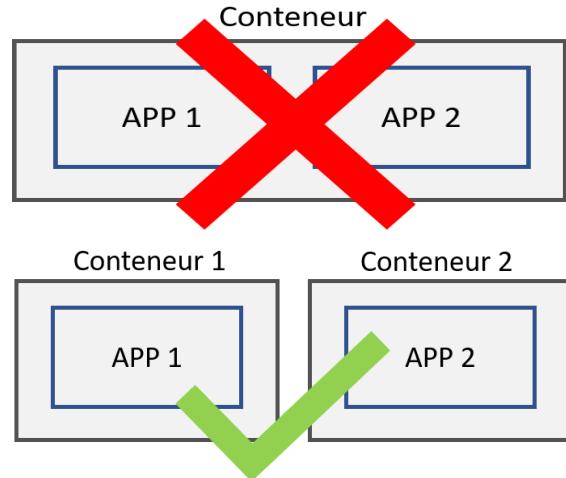


Registres privés

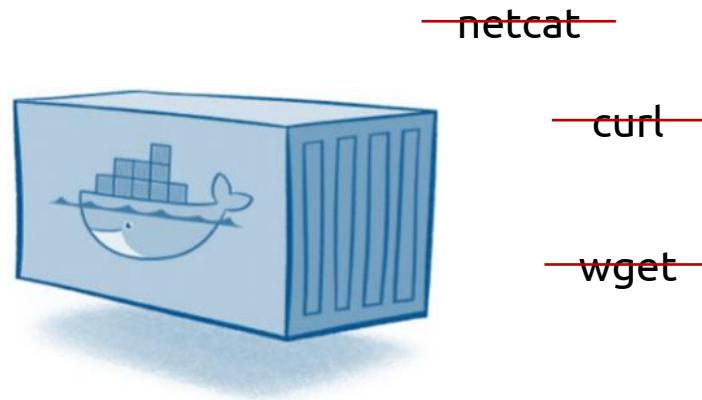


CRÉATION D'UNE IMAGE

Bonne pratique : conteneuriser le minimum



Un seul composant par conteneur
e.g. application et BDD dans des conteneurs différents



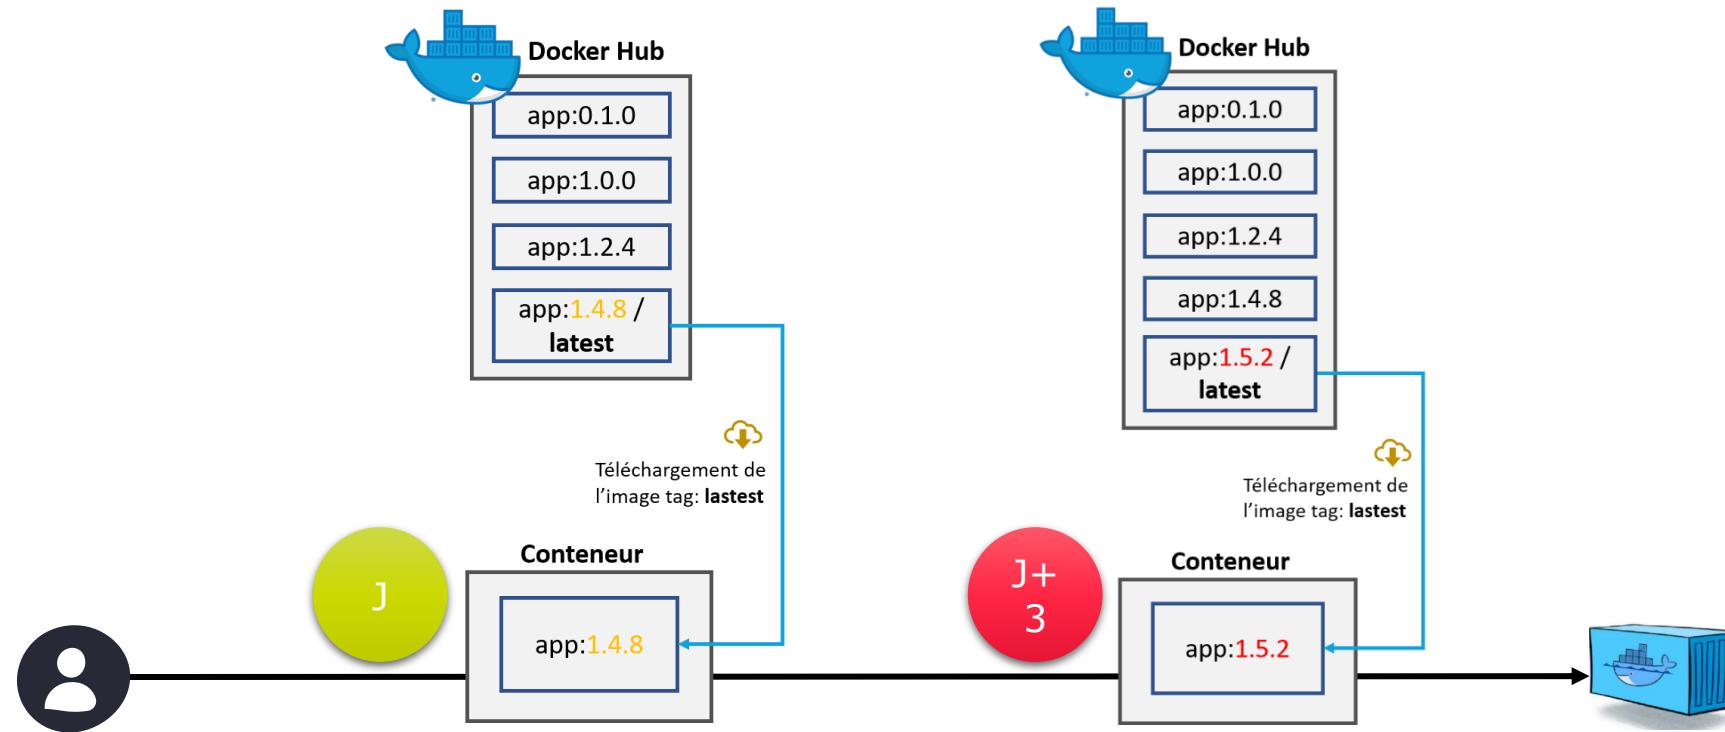
Supprimer les outils inutilisés
e.g. netcat, curl, wget



Réduire poids et temps de démarrage
Réduire la portée d'une action malveillante

CRÉATION D'UNE IMAGE

Bonne pratique : utiliser des images taguées



La version de l'image latest dépend de la date de pull
→ Améliorer la stabilité en réduisant les effets de bords

CRÉATION D'UNE IMAGE

Bonne pratique : utiliser des images taguées

```
FROM ubuntu:latest
RUN apt-get update --fix-missing && \
    apt-get install -y redis-server && \
    rm -rf /var/lib/apt/lists/*
USER 9000
EXPOSE 6379
CMD redis-server
```

✗ Image non taguée

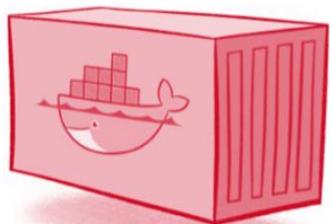


```
FROM ubuntu:xenial
RUN apt-get update --fix-missing && \
    apt-get install -y redis-server && \
    rm -rf /var/lib/apt/lists/*
USER 9000
EXPOSE 6379
CMD redis-server
```

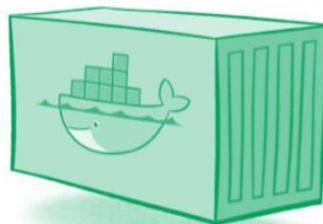
✓ Image taguée

CRÉATION D'UNE IMAGE

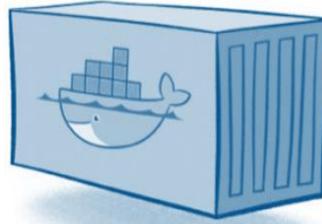
Bonne pratique : utiliser des images non root



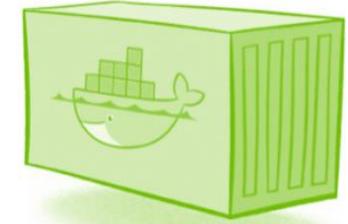
Root



Non-Root



Privileged



Un-privileged

Accès illimité au conteneur
Accès au système de fichiers
de l'hôte (volume)

Accès limité au
conteneur

Accès aux devices
(/dev) de l'hôte

Accès limité aux devices
(/dev) de l'hôte



Construire des images non-root
Capitaliser sur des images non-root existantes
→ Réduire la surface d'attaque

CRÉATION D'UNE IMAGE

Bonne pratique : utiliser des images non root

```
FROM ubuntu:latest
RUN apt-get update --fix-missing && \
    apt-get install -y redis-server && \
    rm -rf /var/lib/apt/lists/*
EXPOSE 6379
CMD redis-server
```

```
$ docker run --rm example whoami
root
```

✗ Utilisateur root (par défaut)



```
FROM ubuntu:xenial
RUN apt-get update --fix-missing && \
    apt-get install -y redis-server && \
    rm -rf /var/lib/apt/lists/*
USER 9000
EXPOSE 6379
CMD redis-server
```

```
$ docker run --rm example whoami
whoami: cannot find name for user ID 9000
```

✓ Utilisateur non root

LANCLEMENT D'UN CONTENEUR

Le conteneur : version dynamique de l'application

Aides

```
docker --help
```

```
docker inspect <id>
```

Listings

```
docker images
```

```
docker ps
```

```
docker rm(i) <id>
```

Inspections

```
docker logs <id>
```

```
docker exec -it <id> <cmd>
```

Instanciations

```
docker build <Dockerfile>
```

```
docker create <img>
```

```
docker run <img>
```



?

?

?



ORCHESTRER SES CONTENEURS VIA KUBERNETES
POUR PASSER À L'ÉCHELLE

Conteneuriser ses applications

- Concepts & bénéfices
- Création d'une image
- Lancement d'un conteneur

Orchestrer ses conteneurs via Kubernetes

- Concepts & bénéfices
- Déployer un cluster Kubernetes
- Les ressources Kubernetes
- Déployer des ressources sur le cluster

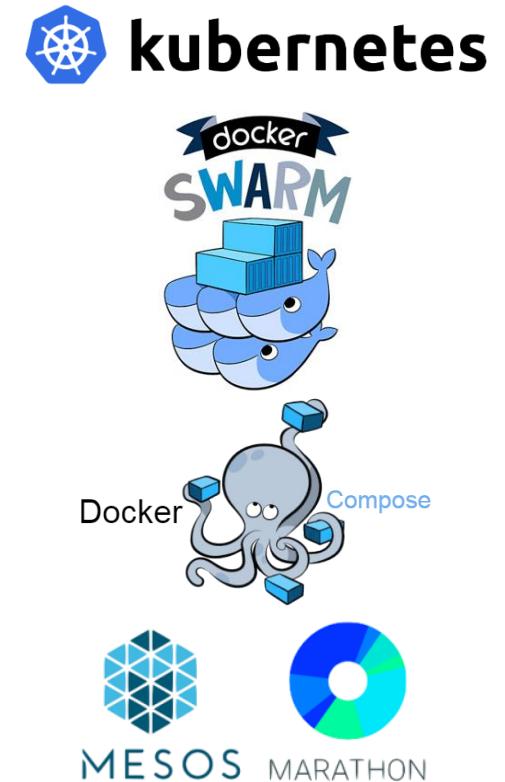
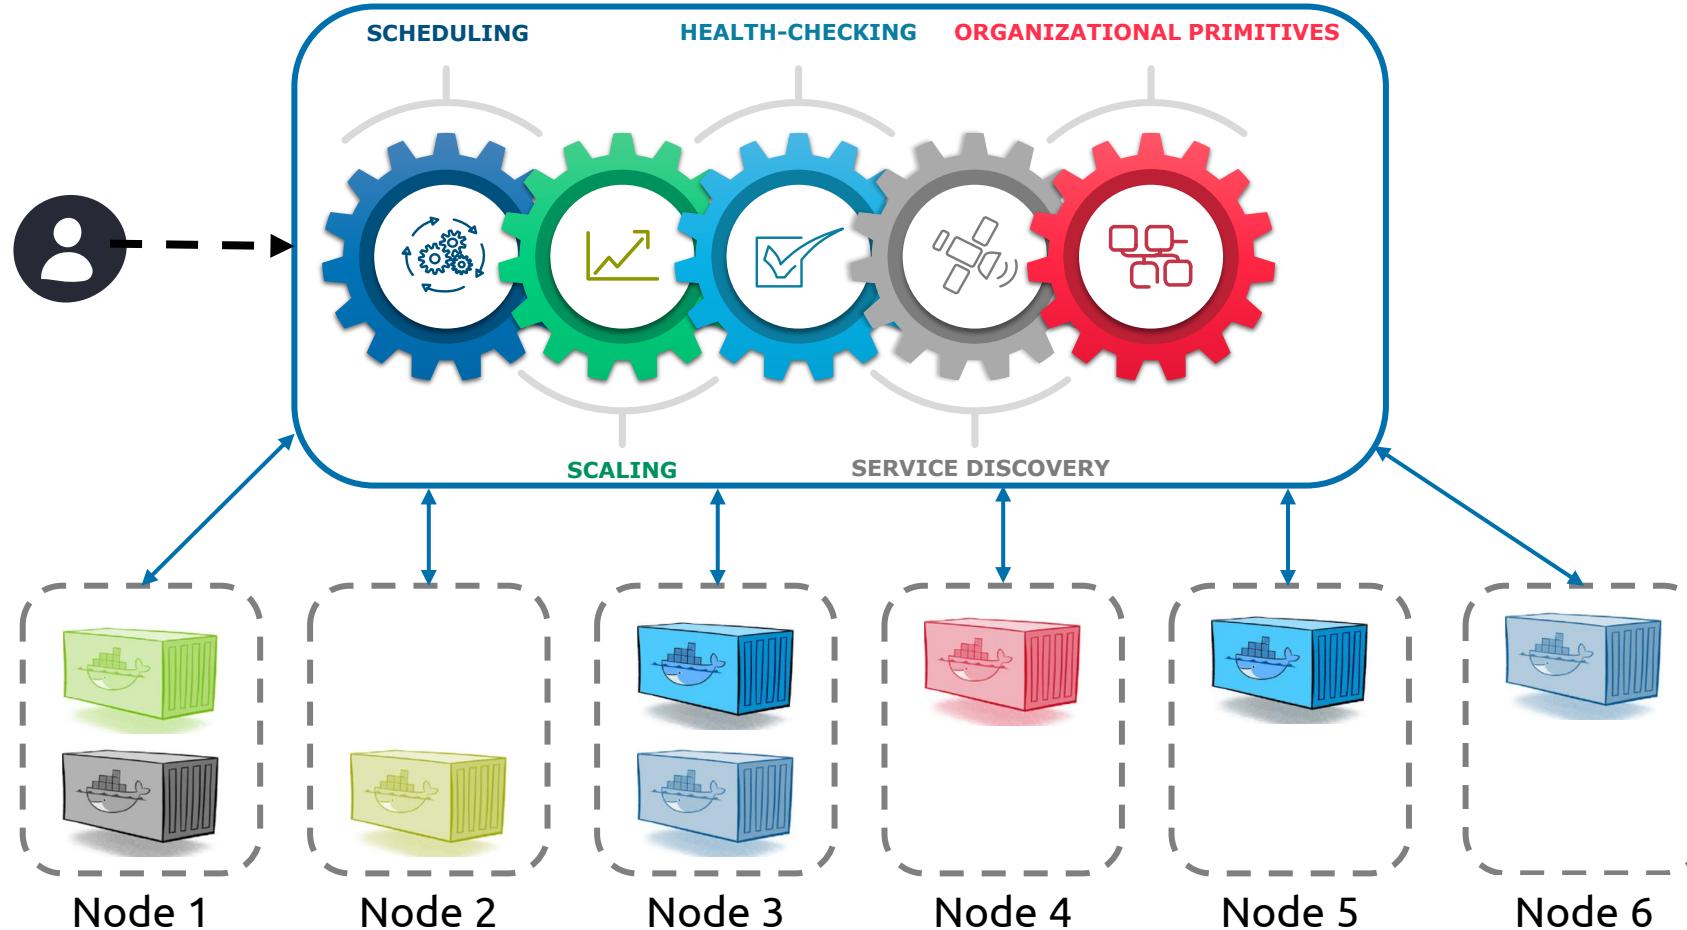
Industrialiser ses déploiements

- Du code au cluster
- Automatiser la construction et le déploiement

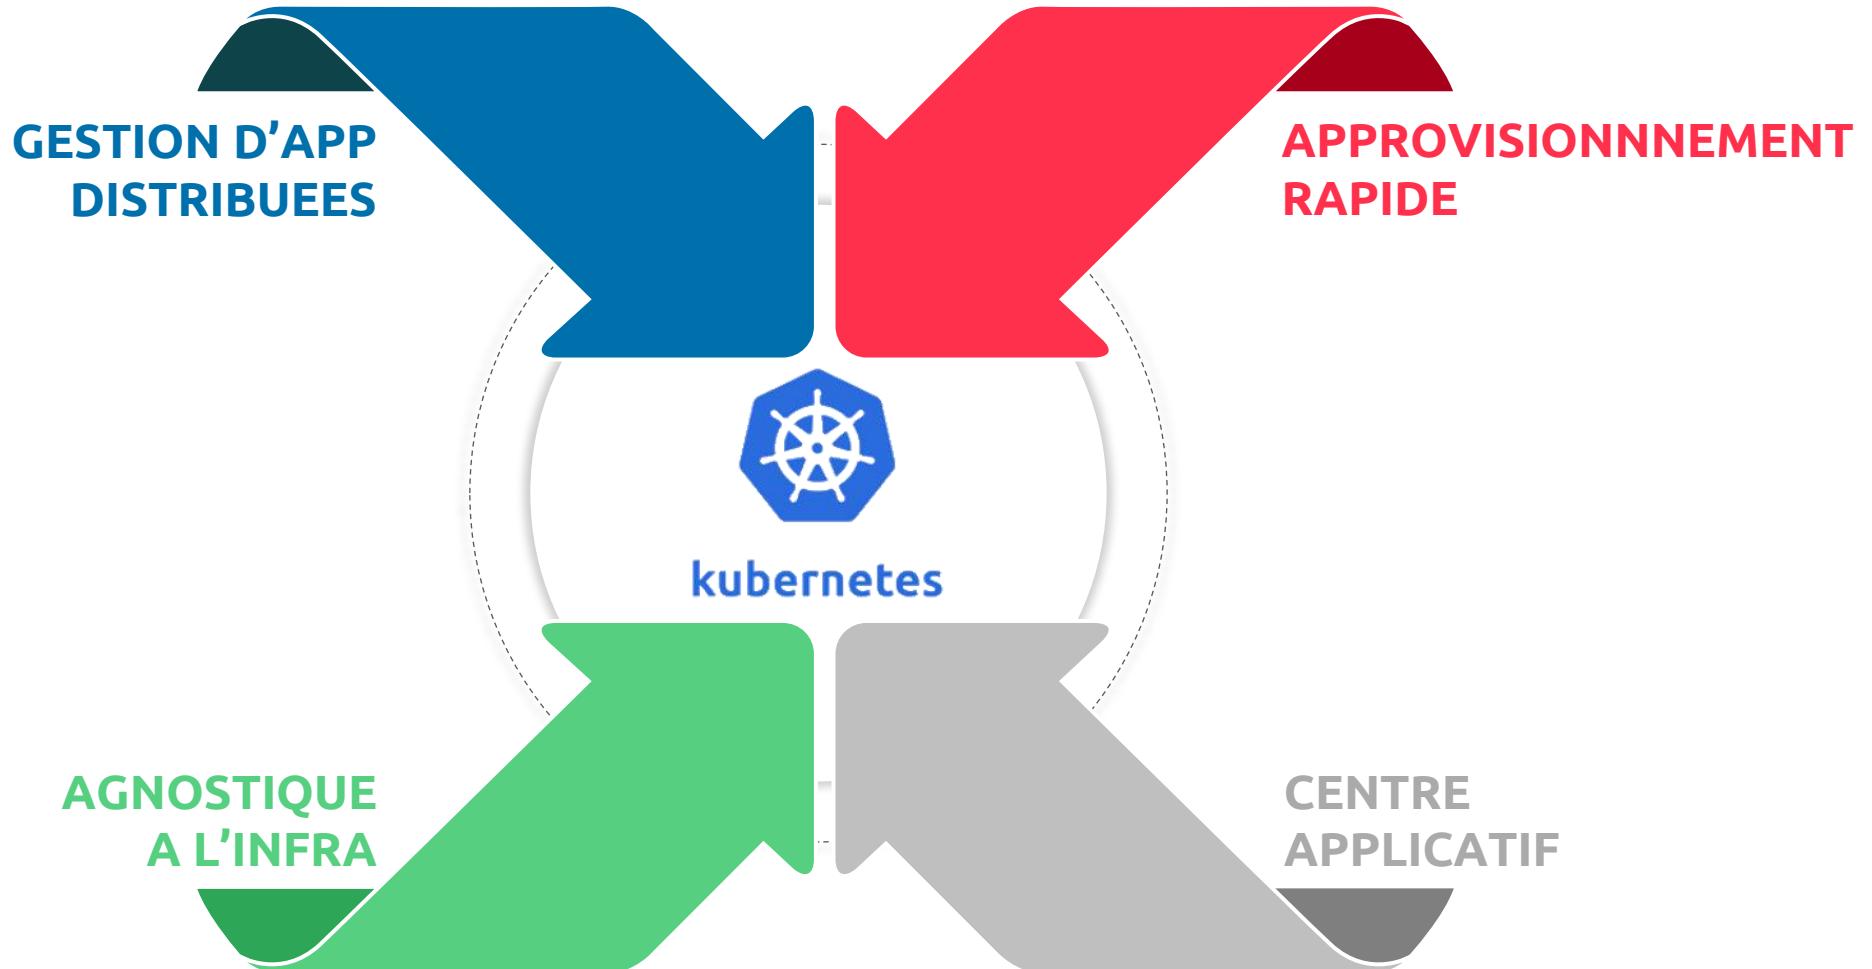
Exploiter & superviser ses applications en production

- Exploiter ses applications
- Superviser ses applications
- Collecter des métriques

DÉPLOYER DES CONTENEURS SUR DES MACHINES

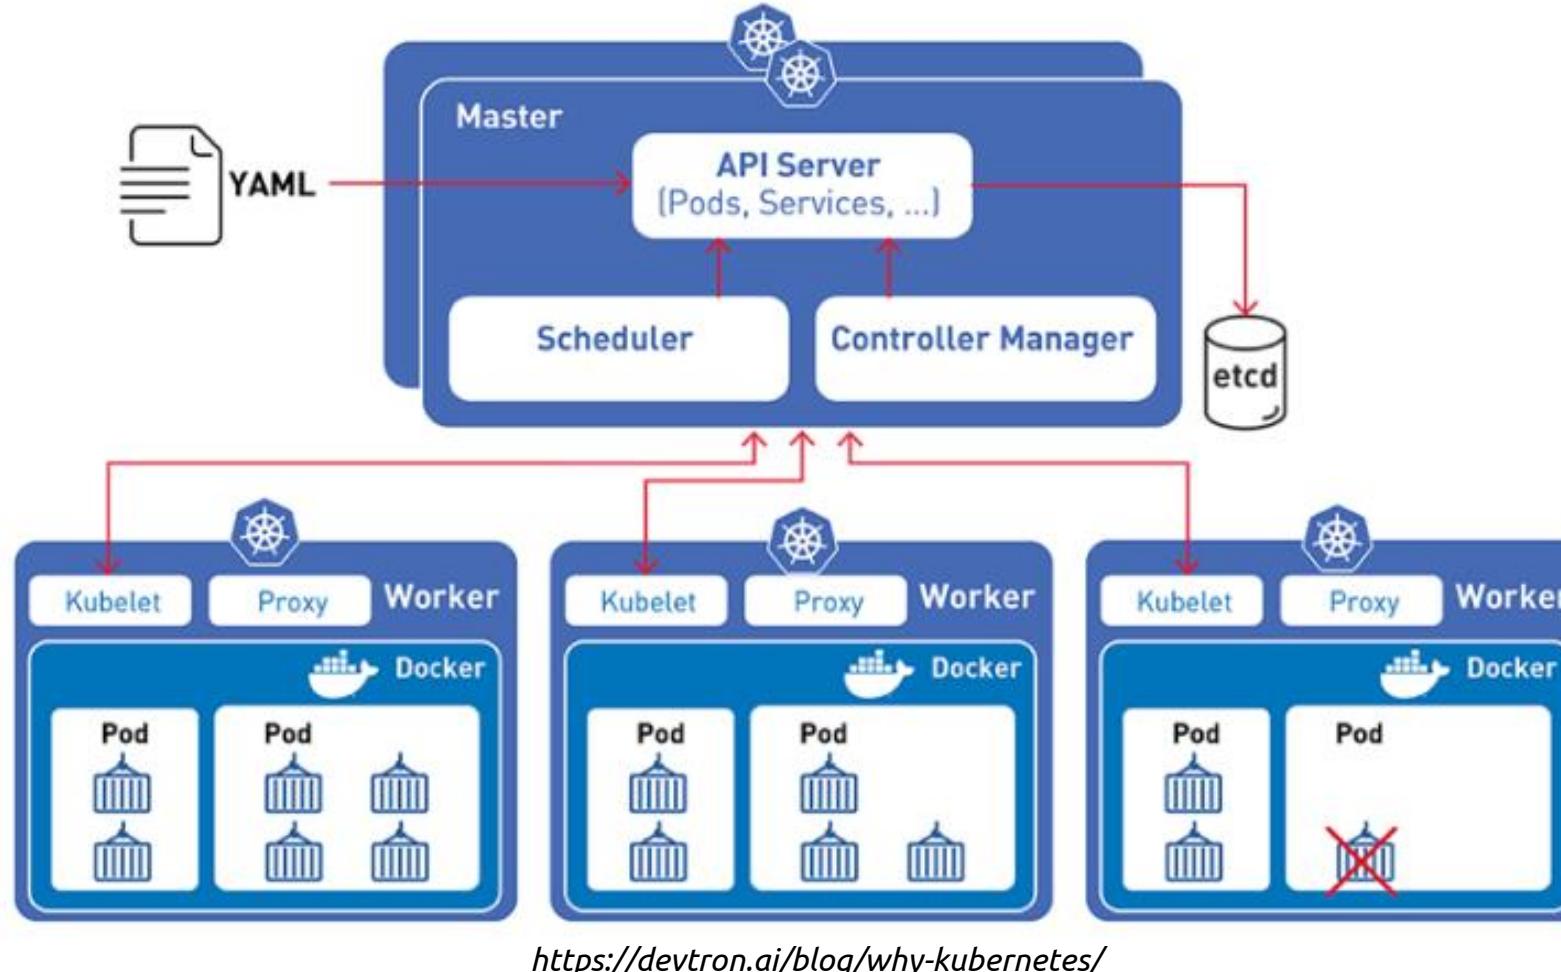


BÉNÉFICES DE L'ORCHESTRATION



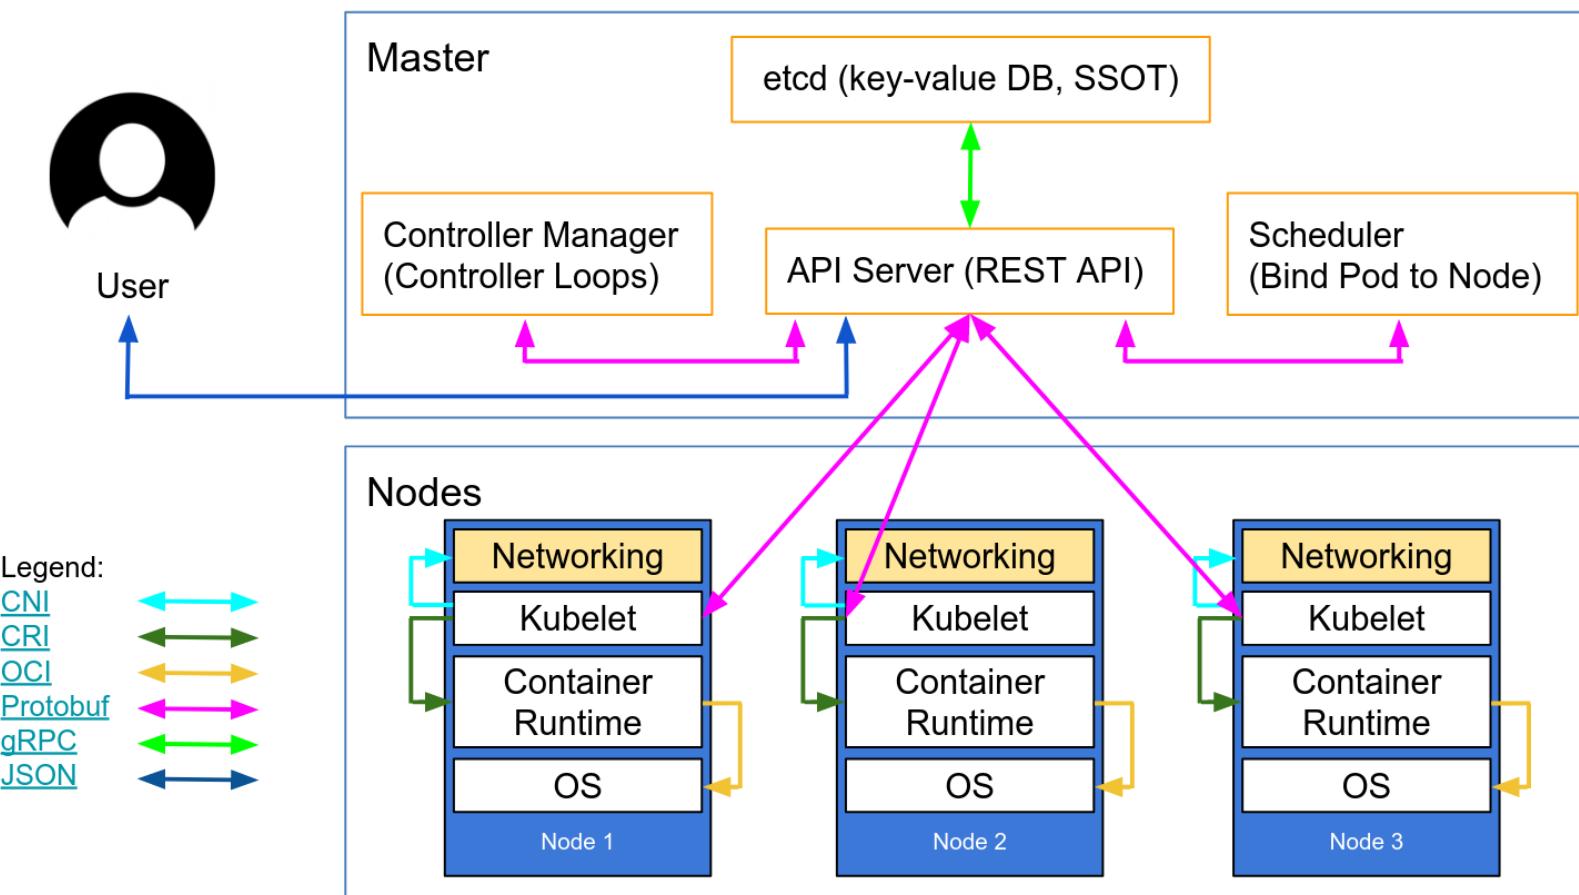
ARCHITECTURE D'UN CLUSTER KUBERNETES

Kubernetes, un orchestrateur modulaire



ARCHITECTURE D'UN CLUSTER KUBERNETES

Flux d'informations au sein d'un cluster



Composants Master

- kube-apiserver
- kube-scheduler
- etcd
- controller-manager

Composants nœuds

- kubelet
- kube-proxy ou CN (CNI)
- CR (CRI)

DÉPLOYER UN CLUSTER KUBERNETES

Choix de l'installer / manager



Installers

Kops / Kubespray

- Déploie l'infra & installe K8s
- Nombre limité de providers
- Debug difficile



Cloud-Hosted K8 Services

GKE / EKS / AKS

- Facilité de mise en place
- Composants K8s managés



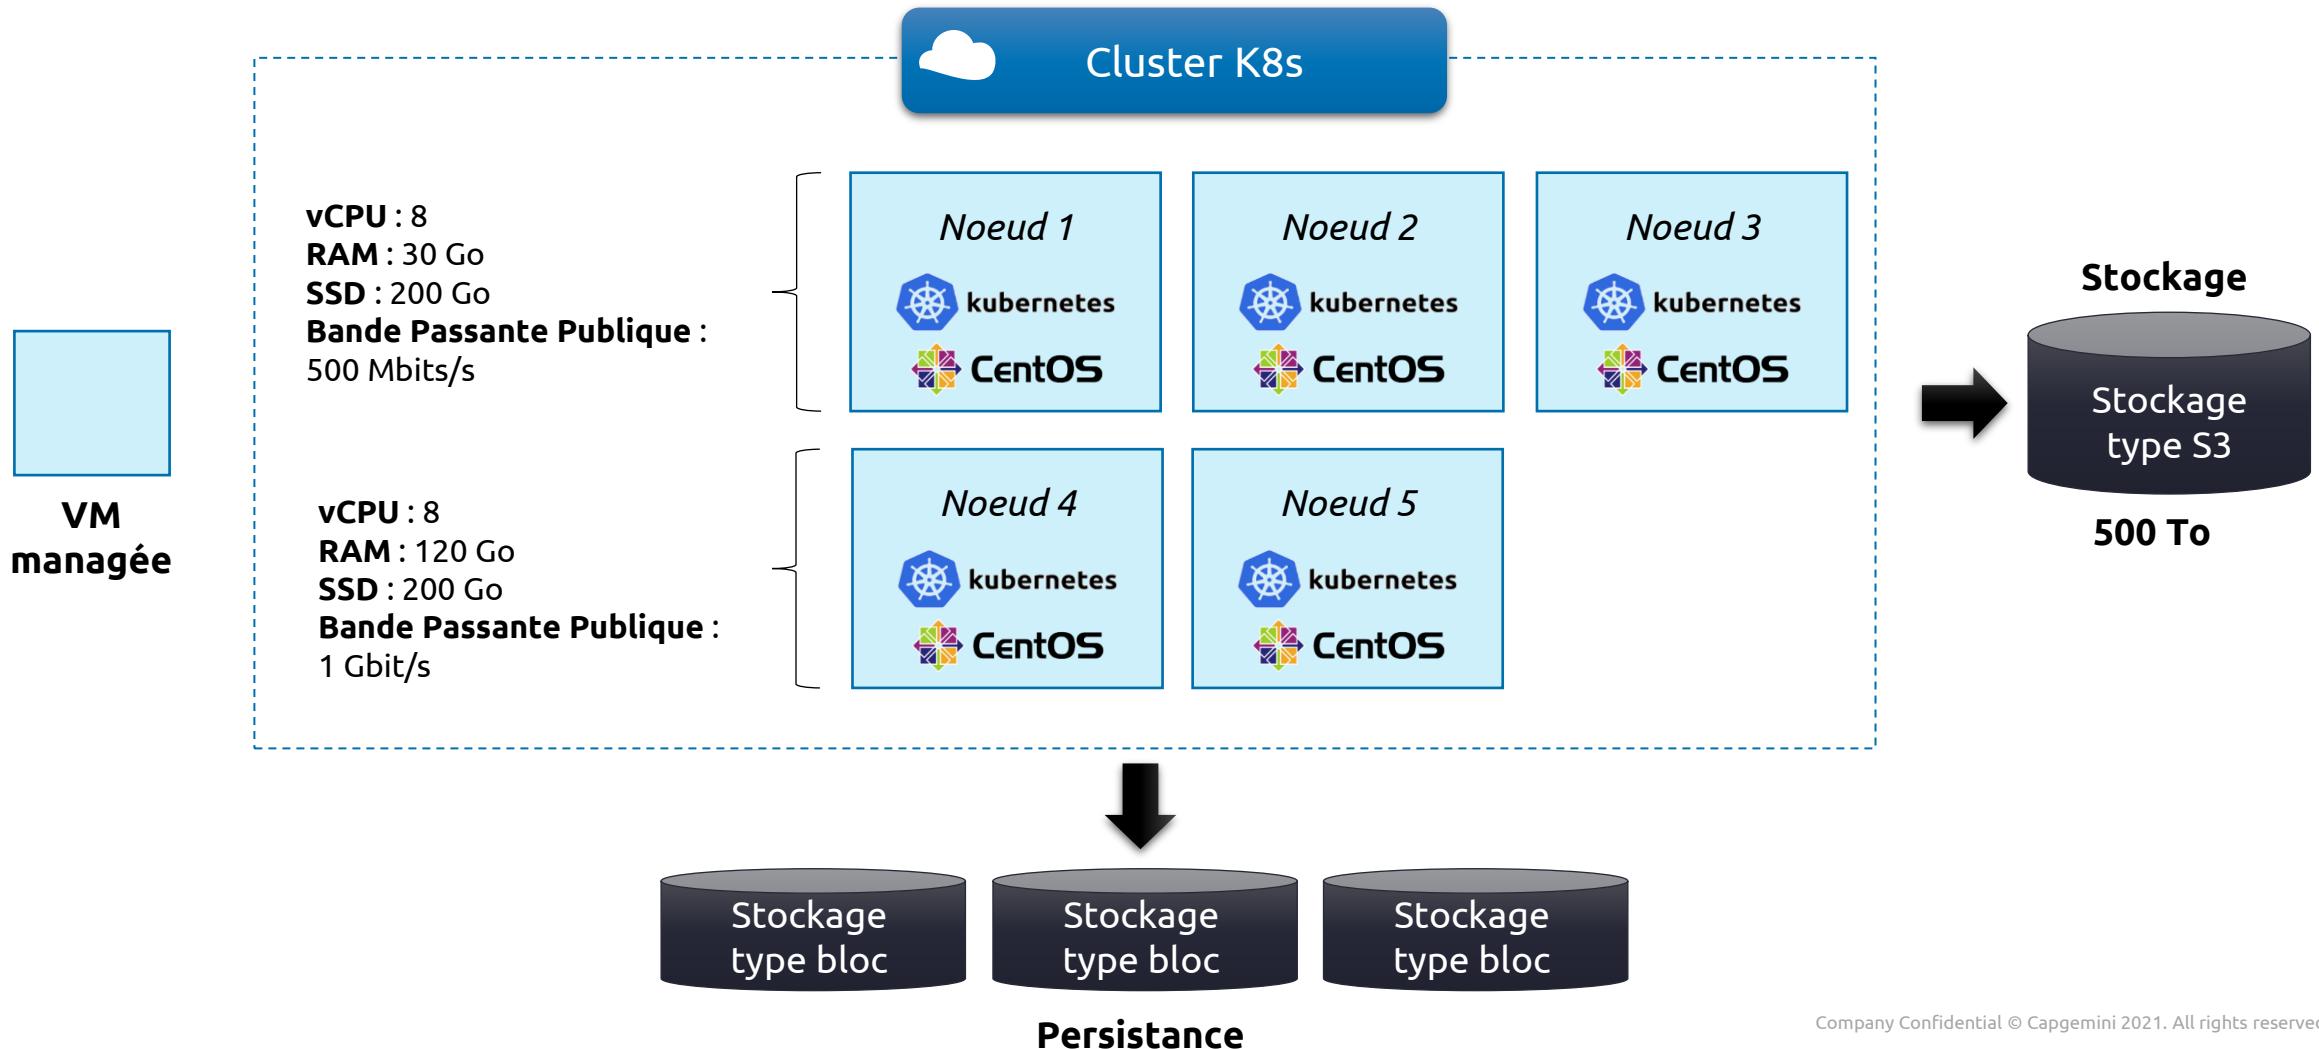
Red Hat Kubernetes Installer

Red Hat OpenShift

- Installeur automatisé
- Fonctionnalités additionnelles
- Payant

DÉPLOYER UN CLUSTER KUBERNETES

Exemple d'architecture physique





?

?

?

Conteneuriser ses applications

- Concepts & bénéfices
- Création d'une image
- Lancement d'un conteneur

Orchestrer ses conteneurs via Kubernetes

- Concepts & bénéfices
- Déployer un cluster Kubernetes
- Les ressources Kubernetes
- Déployer des ressources sur le cluster

Industrialiser ses déploiements

- Du code au cluster
- Automatiser la construction et le déploiement

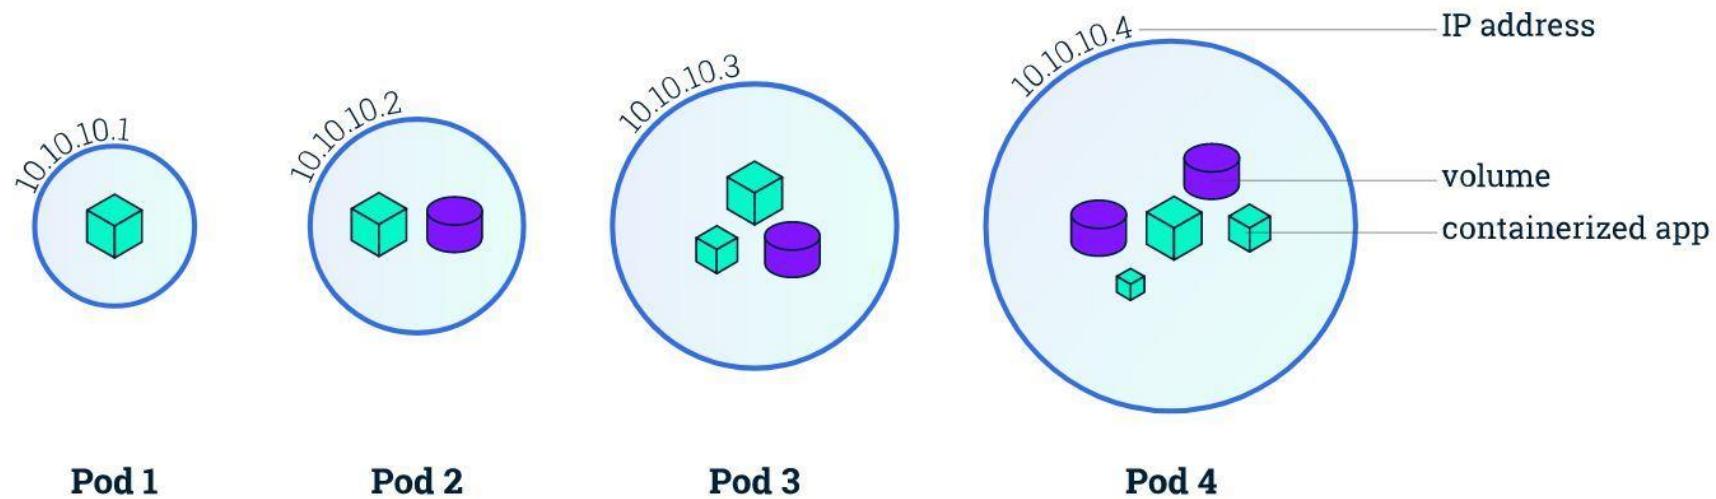
Exploiter & superviser ses applications en production

- Exploiter ses applications
- Superviser ses applications
- Collecter des métriques

POD

Unité élémentaire de Kubernetes

Ressource Kubernetes : objet déployé par un *manifest yaml*



DEPLOYMENT

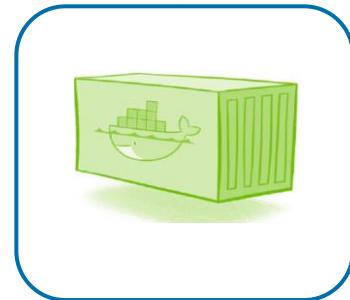
Déployer son application sur un cluster Kubernetes



deployment.yaml

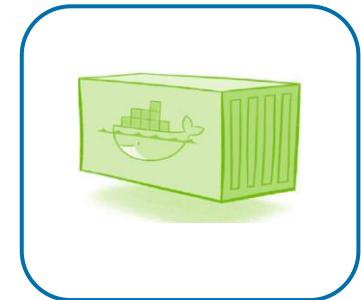
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-petstore
  namespace: java
spec:
  selector:
    matchLabels:
      run: my-petstore
  replicas: 2
  template:
    metadata:
      labels:
        run: my-petstore
    spec:
      containers:
        - name: petstore
          image: <registre>/java-petstore-test:latest
```

Pod : my-petstore-1



10.10.10.1

Pod : my-petstore-2

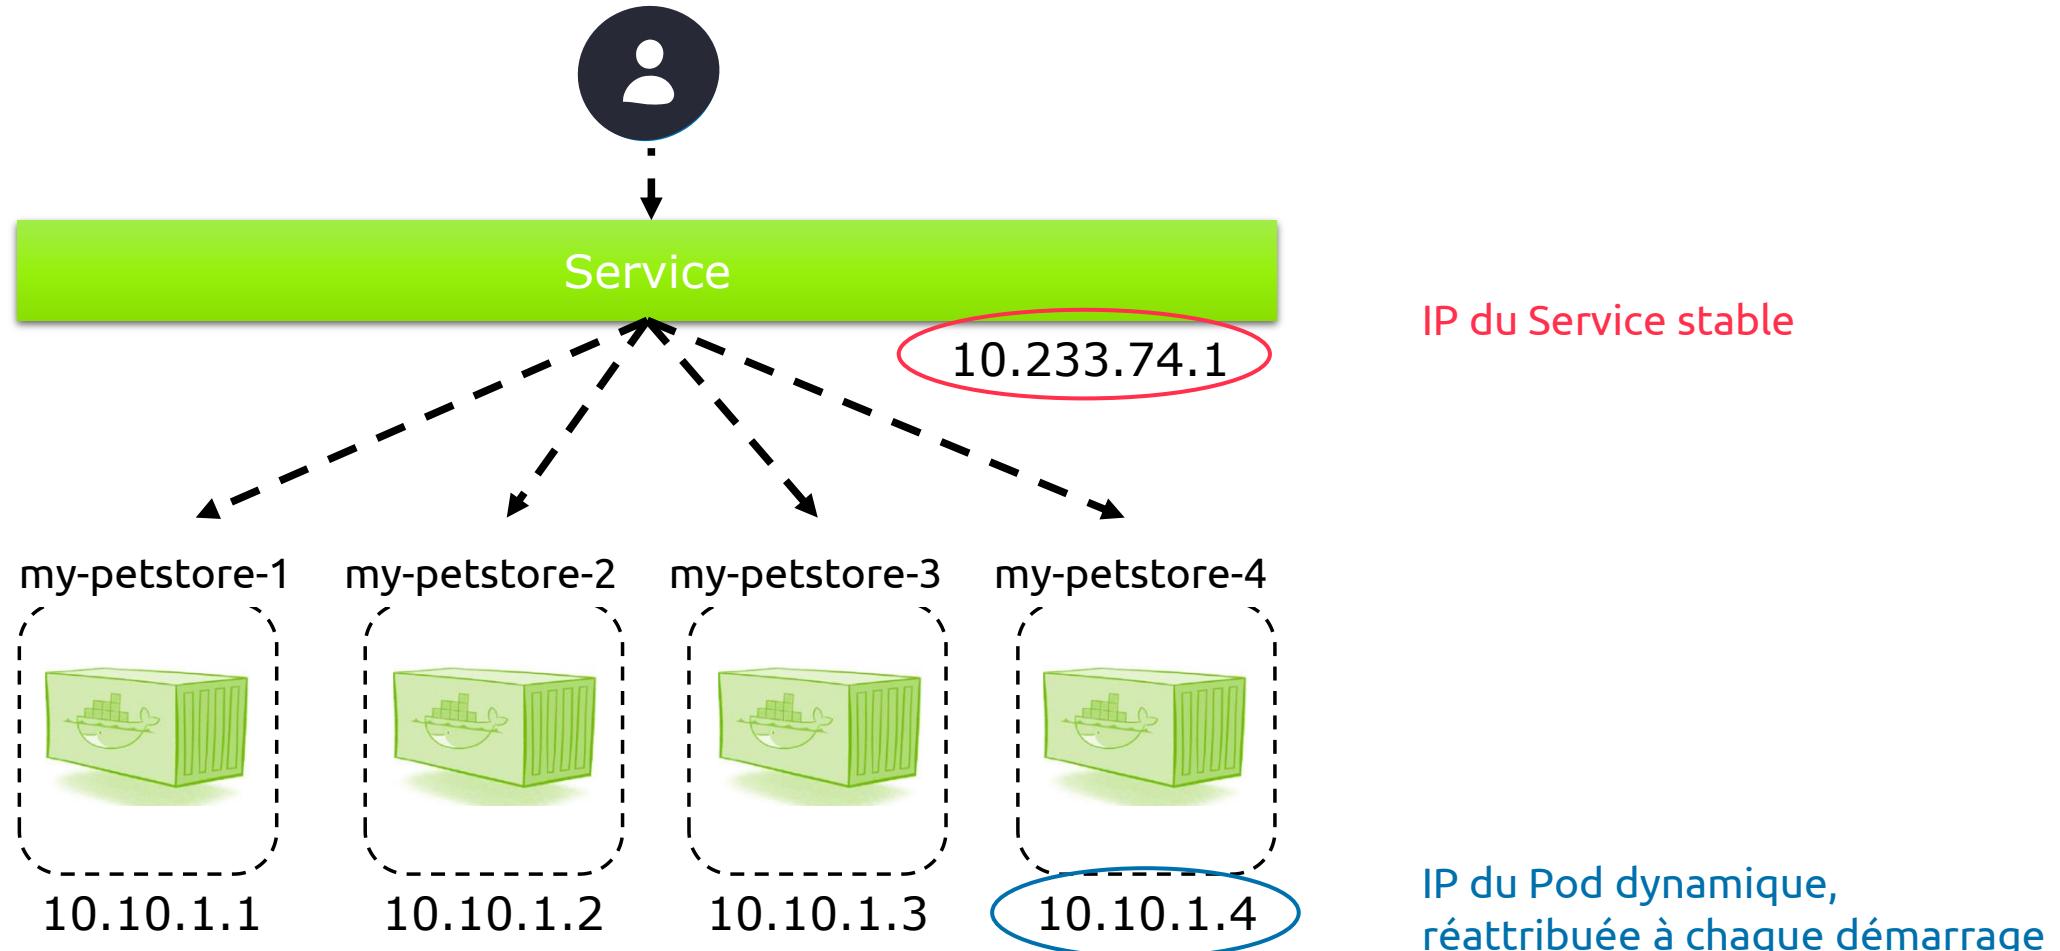


10.10.10.2

- Stabilité de l'application
- Continuité du service
- Facilité de déploiement

SERVICE

Loadbalancer intégré à Kubernetes



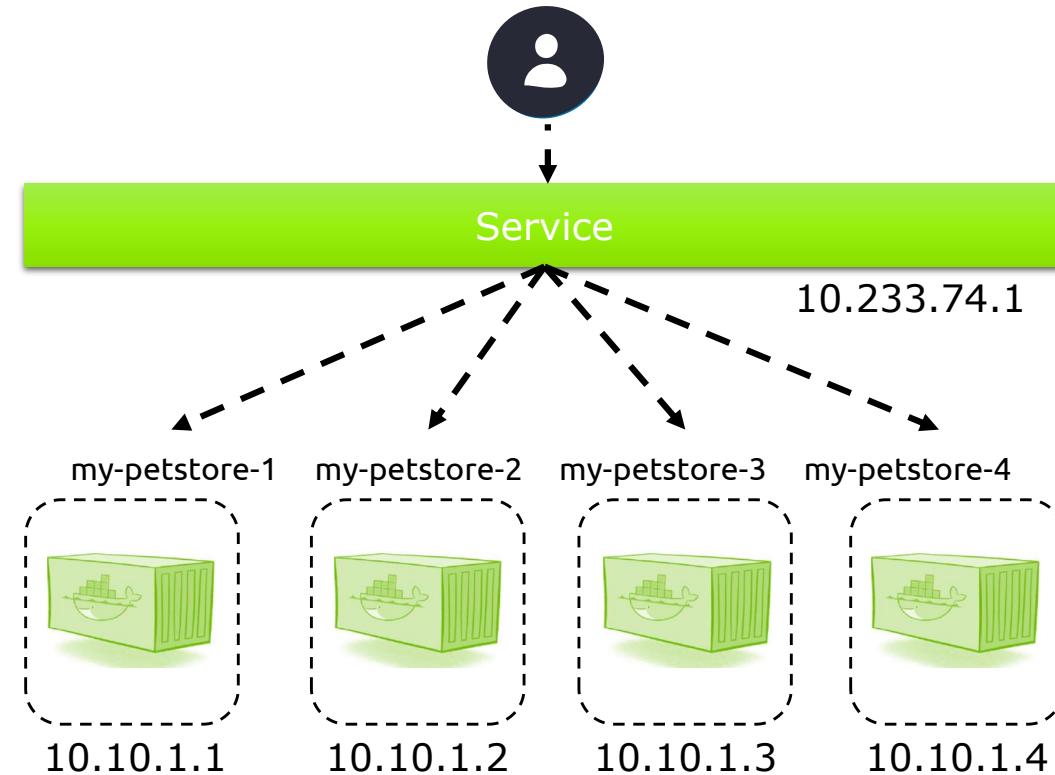
SERVICE

Déployer un Service



service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: my-petstore
  namespace: java
  labels:
    run: my-petstore
spec:
  ports:
    - port: 8080
      protocol: TCP
  selector:
    run: my-petstore
```



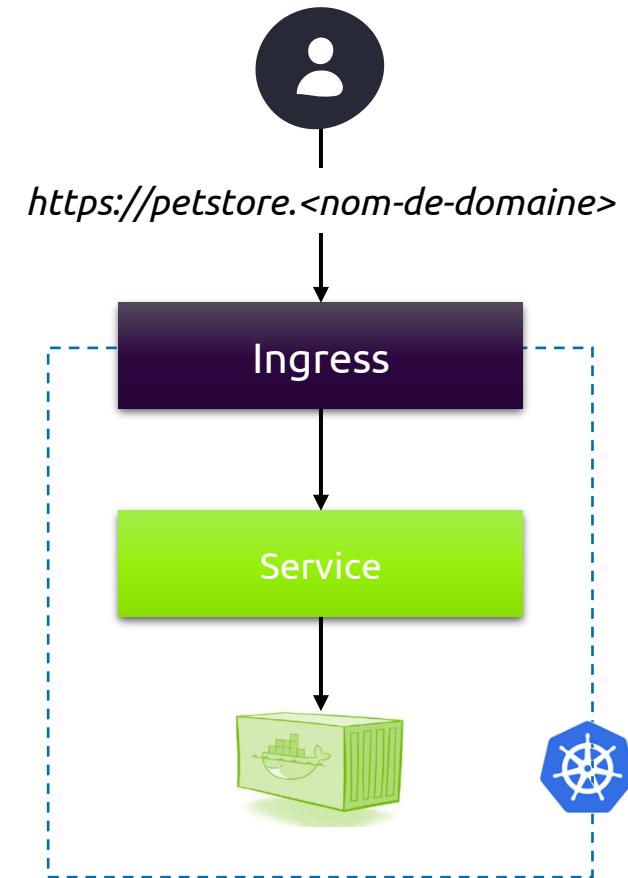
INGRESS

Configurer le trafic entrant avec l'Ingress



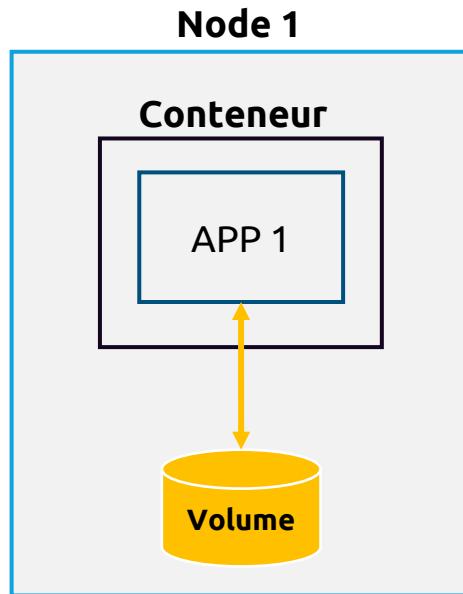
ingress.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/proxy-connect-timeout: "30"
    nginx.ingress.kubernetes.io/ssl-redirect: "false"
    nginx.ingress.kubernetes.io/app-root: /applicationPetstore
  labels:
    app: my-petstore
    name: my-petstore
    namespace: java
spec:
  rules:
    - host: petstore. <nom-de-domaine>
      http:
        paths:
          - backend:
              serviceName: my-petstore
              servicePort: 8080
```



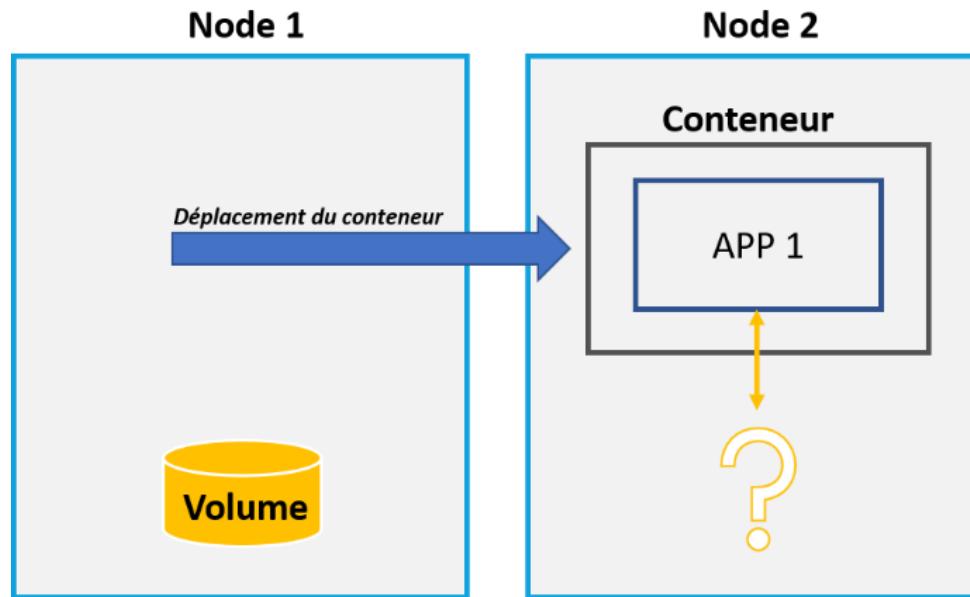
VOLUMES

Persister la donnée sur un stockage bloc

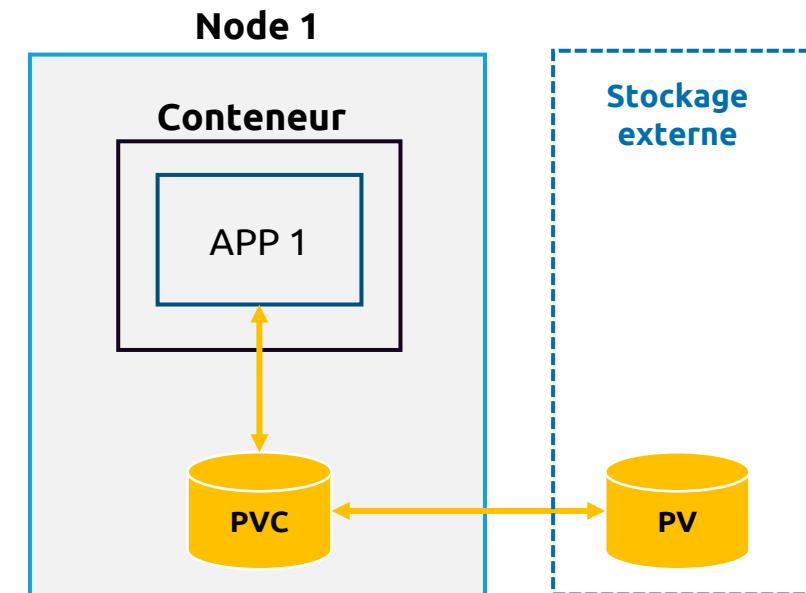


VOLUMES

Persister la donnée sur un stockage bloc externe



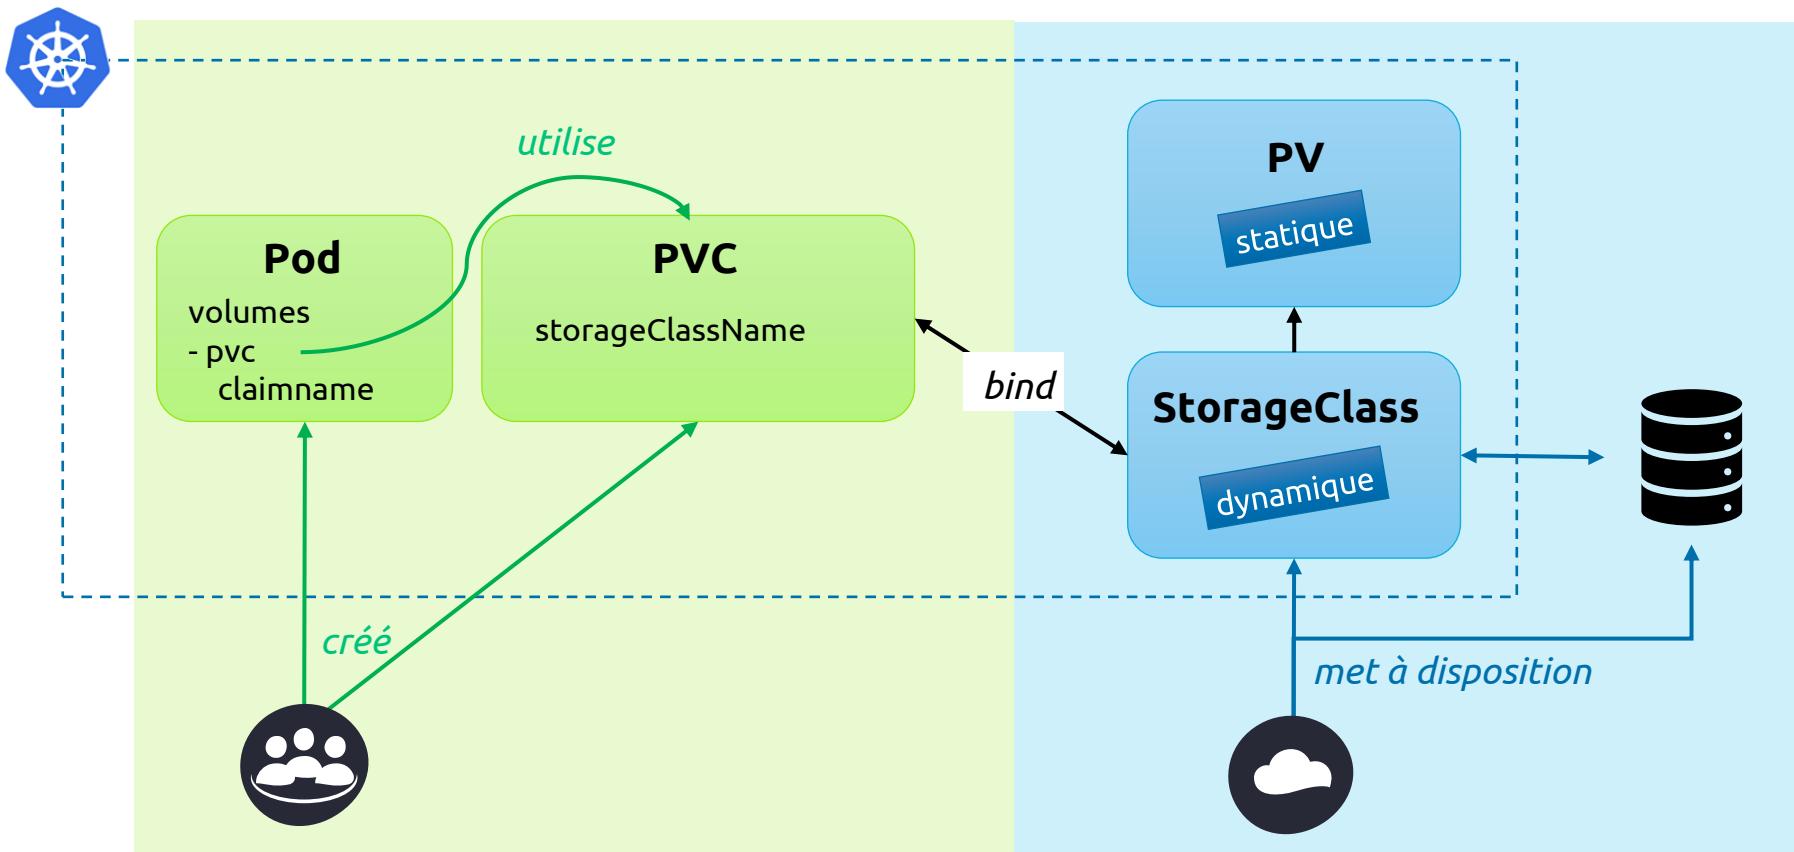
✗ **Volume intra cluster**



✓ **Volume externe**

VOLUMES

PersistentVolume & PersistentVolumeClaim



VOLUMES

PersistentVolumeClaim : demander un volume via une storageClass



pvc.yaml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: bdd-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: csi-cinder-high-speed
```

Utiliser une storageClass dynamique
→ Pas de gestion des volumes
→ Approvisionnement dynamique



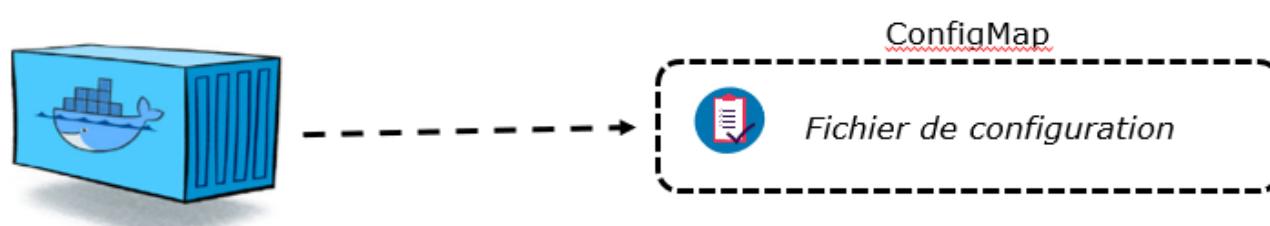
Stabilité de l'application

CONFIGMAPS & SECRETS

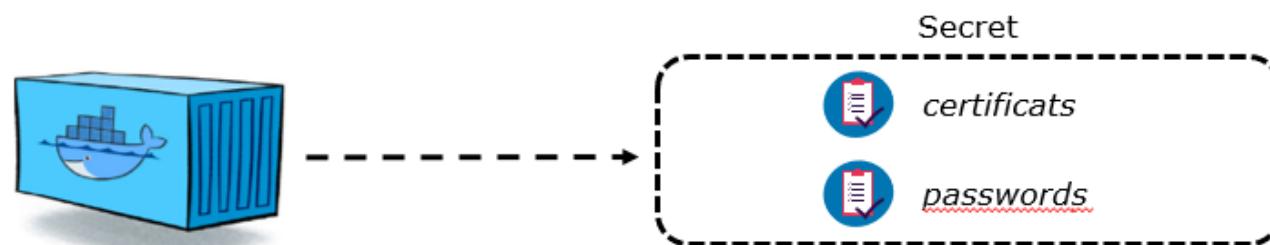
Stocker la configuration et les credentials dans les Pods



Stocker des configurations avec la ConfigMap



Stocker les credentials ou les certificats avec les Secrets



CONFIGMAP

Stocker la configuration des applications dans les Pods



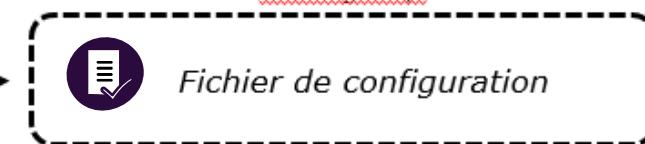
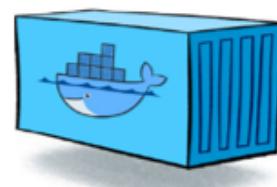
configmap.yaml

```
apiVersion: v1
kind: ConfigMap
data:
  demo: |
    hello world
metadata:
  name: demo
  namespace: java
```



deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
spec:
  template:
    spec:
      volumes:
        - name: demo
          configMap:
            name: demo
```



SECRET

Stocker les secret encodés des applications dans les Pods



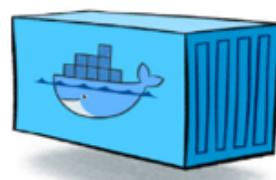
secret.yaml

```
apiVersion: v1
kind: Secret
data:
  username: YWRtaW4=
metadata:
  name: demo
```



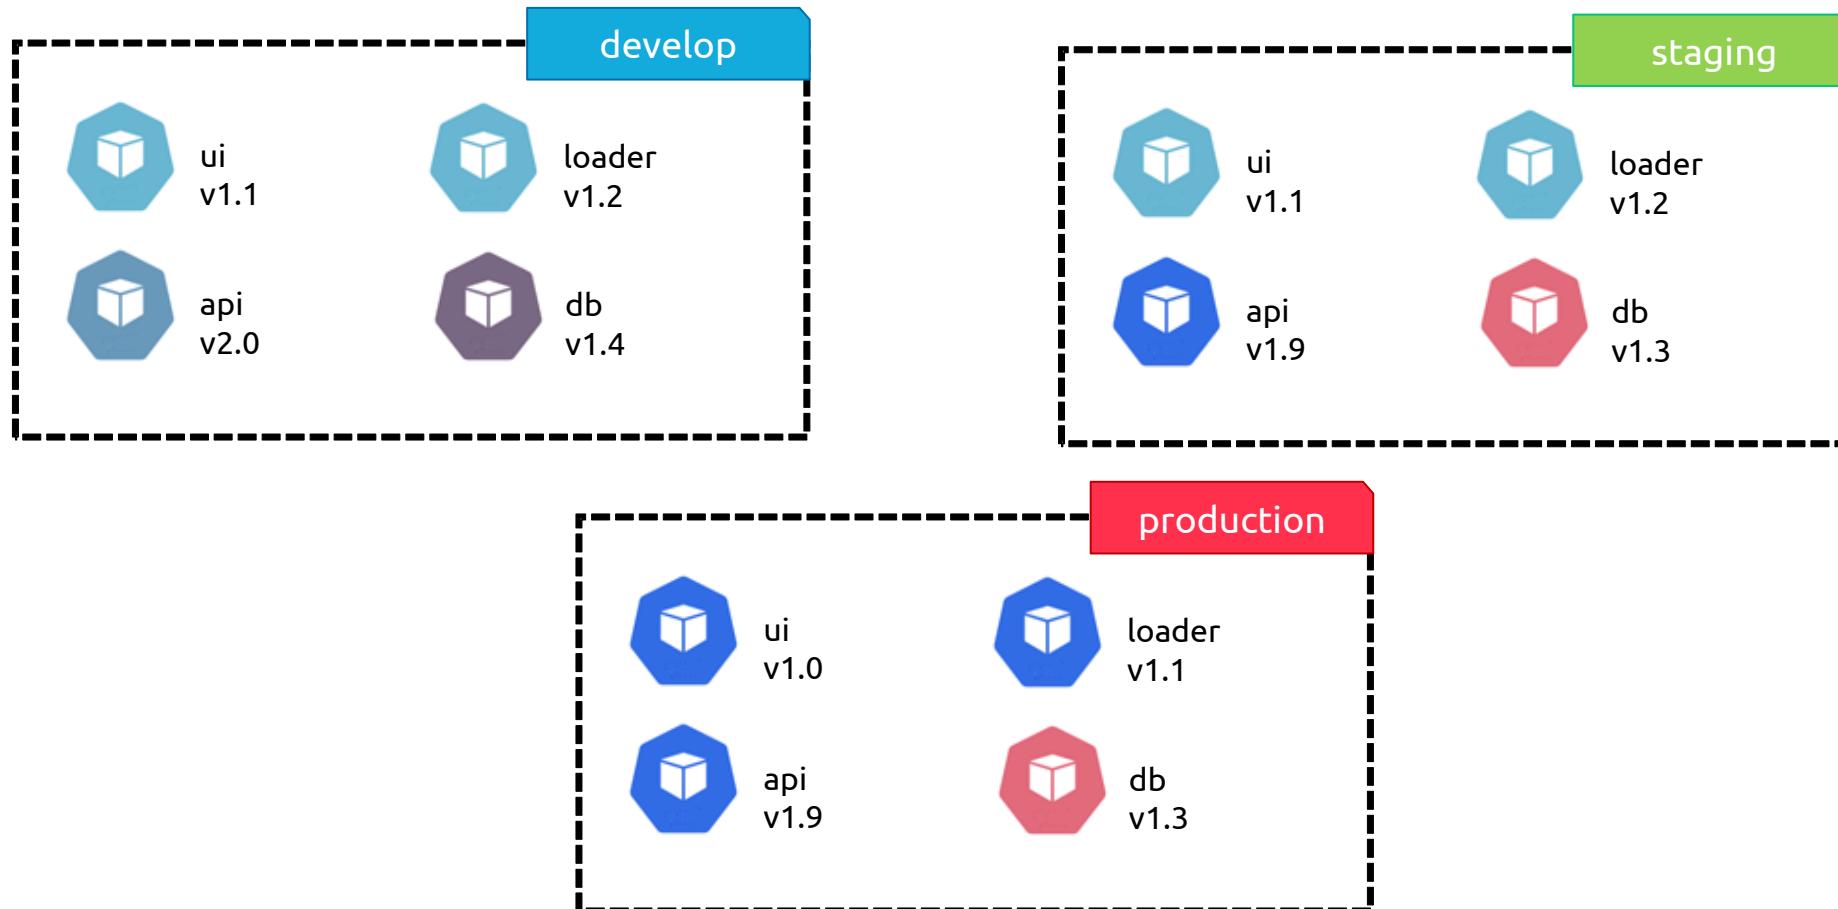
deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
spec:
  template:
    spec:
      volumes:
        - name: demo
          secret:
            secretname: demo
```



NAMESPACE

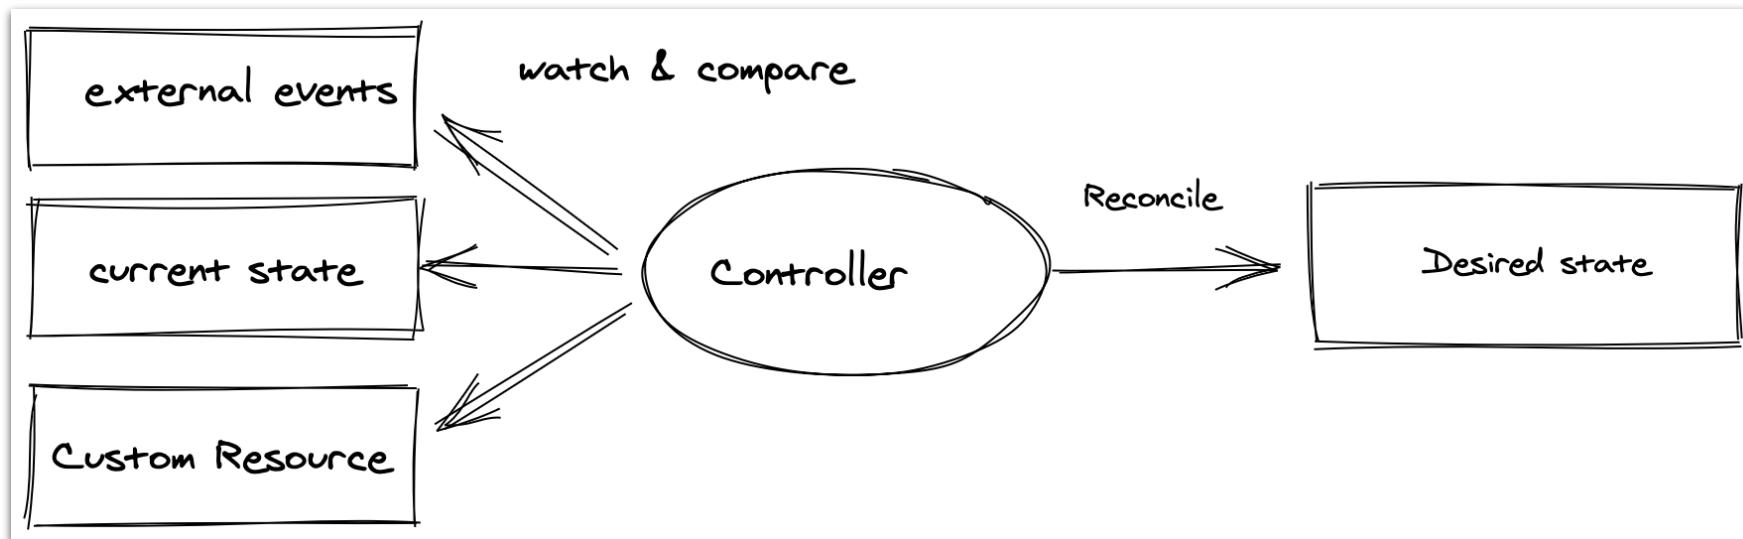
Isoler les Pods via des Namespaces



CUSTOM RESOURCE DEFINITIONS

Extension de l'API Kubernetes par les Operators

"An operator is a Kubernetes controller that understands 2 domains: Kubernetes and something else. By combining knowledge of both domains, it can automate tasks that usually require a human operator that understands both domains." Jimmy Zelinskie



https://github.com/cncf/tag-app-delivery/blob/eece8f7307f2970f46f100f51932db106db46968/operator-wg/whitepaper/Operator-WhitePaper_v1-0.md

CUSTOM RESOURCE DEFINITIONS

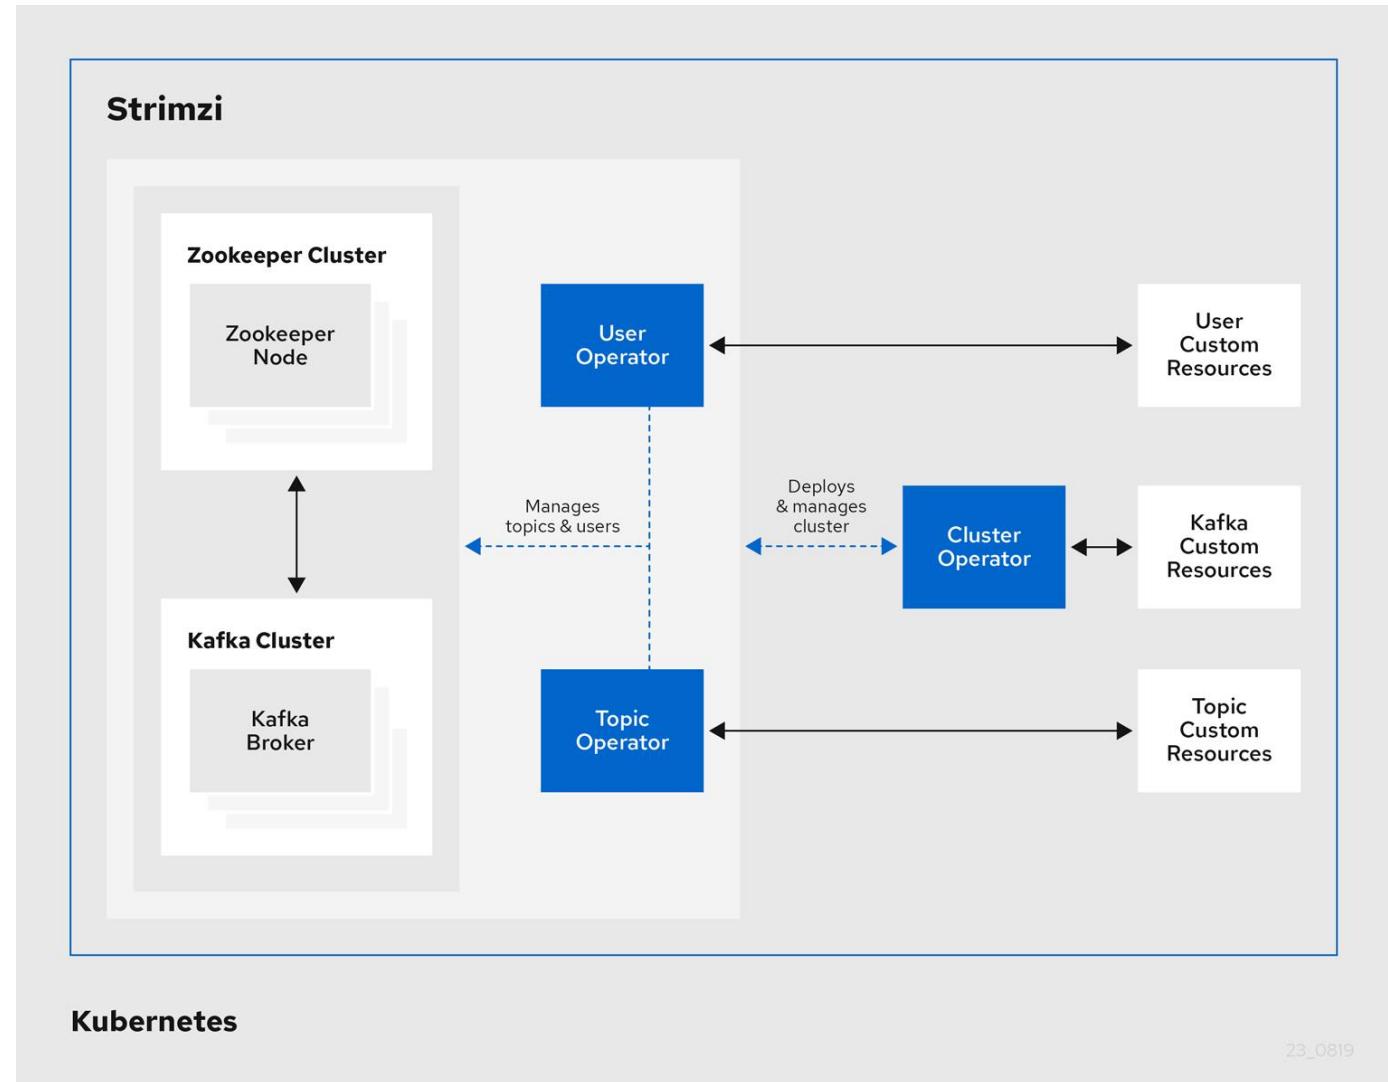
Exemples de Kubernetes Operators



kafka



STRIMZI



DÉPLOYER UN WORKLOAD

CLI kubectl : commandes utiles

Aides

```
kubectl --help
```

```
kubectl explain <resource>
```

Listings

```
kubectl get nodes
```

```
kubectl get pods
```

```
kubectl get svc
```

Inspections

```
kubectl describe pod <pod-id>
```

```
kubectl logs <pod-id>
```

Instanciations

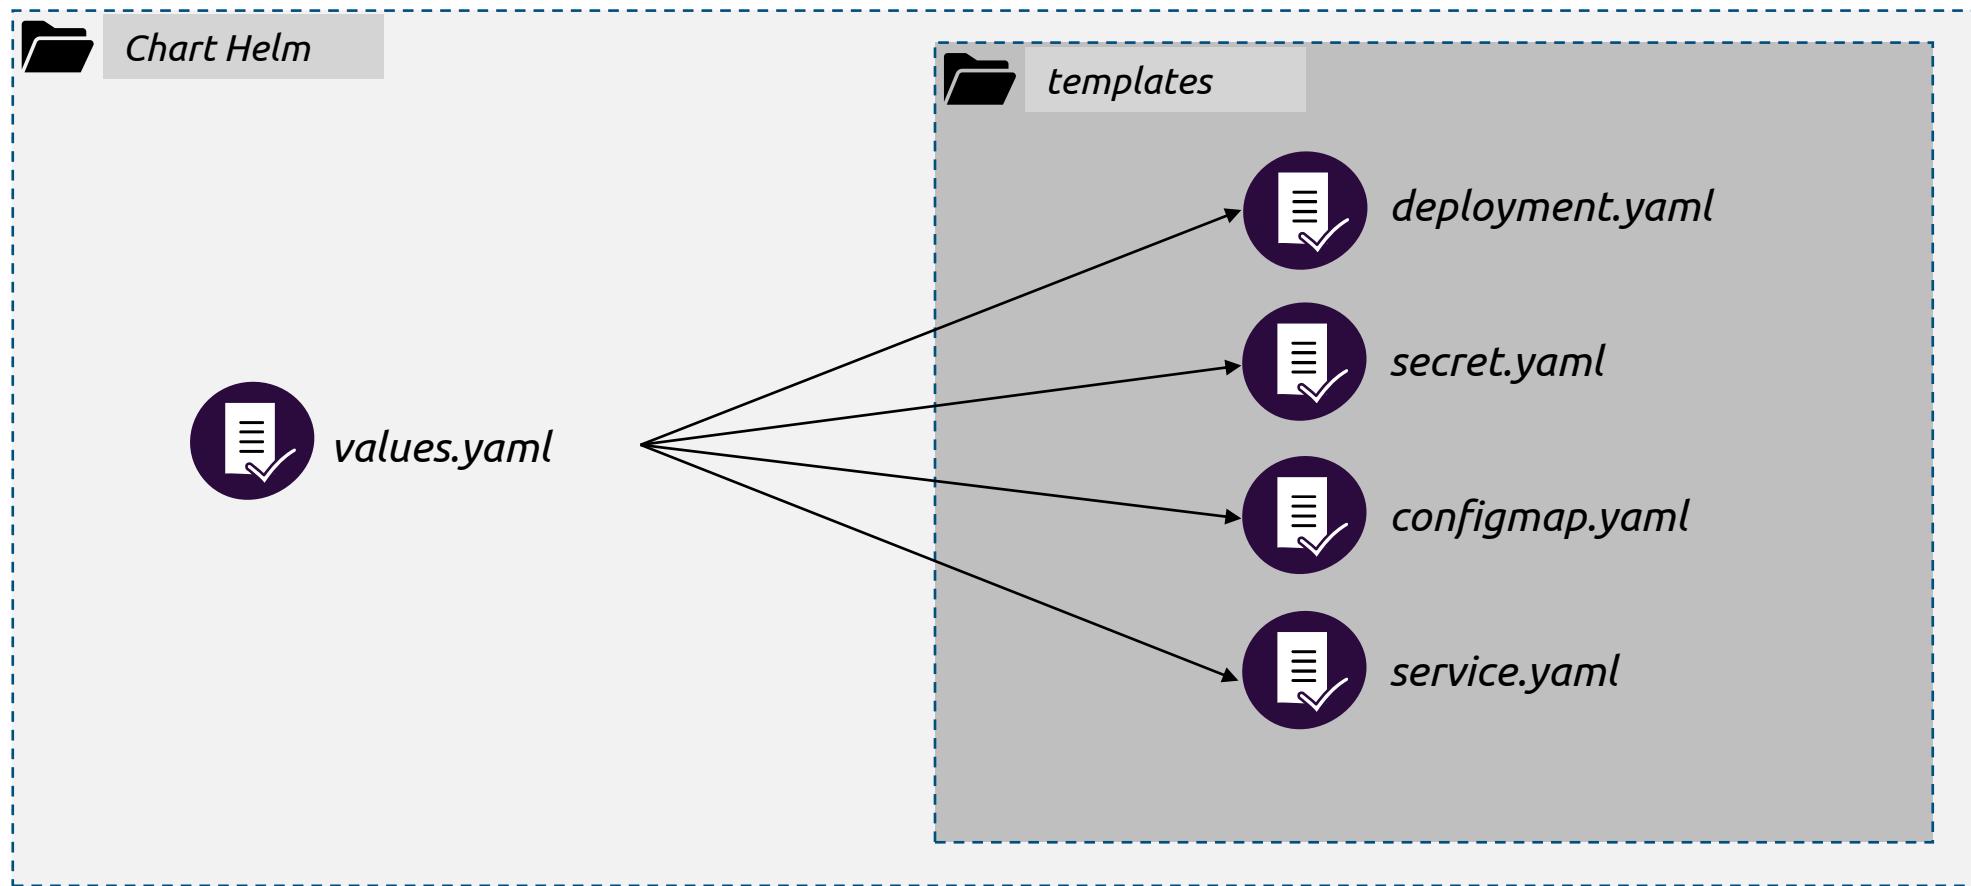
```
kubectl create -f <file>
```

```
kubectl apply -f <file>
```

```
kubectl delete -f <file>
```

DÉPLOYER UN WORKLOAD

Via le manager de paquet Helm & des ressources variabilisées



DÉPLOYER UN WORKLOAD

Commandes utiles de l'utilitaire helm

Aides

```
helm --help
```

Listings

```
helm ls
```

Installation

```
helm install <chart> --name <nom-chart> --namespace <nom-namespace>
```

Télécharger la Helm Chart

```
helm fetch <chart>
```

Désinstallation

```
helm uninstall <nom-chart>
```



?

?

?



DU CODE AU CLUSTER : INDUSTRIALISER SES DÉPLOIEMENTS



Conteneuriser ses applications

- Concepts & bénéfices
- Création d'une image
- Lancement d'un conteneur

Orchestrer ses conteneurs via Kubernetes

- Concepts & bénéfices
- Déployer un cluster Kubernetes
- Les ressources Kubernetes
- Déployer des ressources sur le cluster

Industrialiser ses déploiements

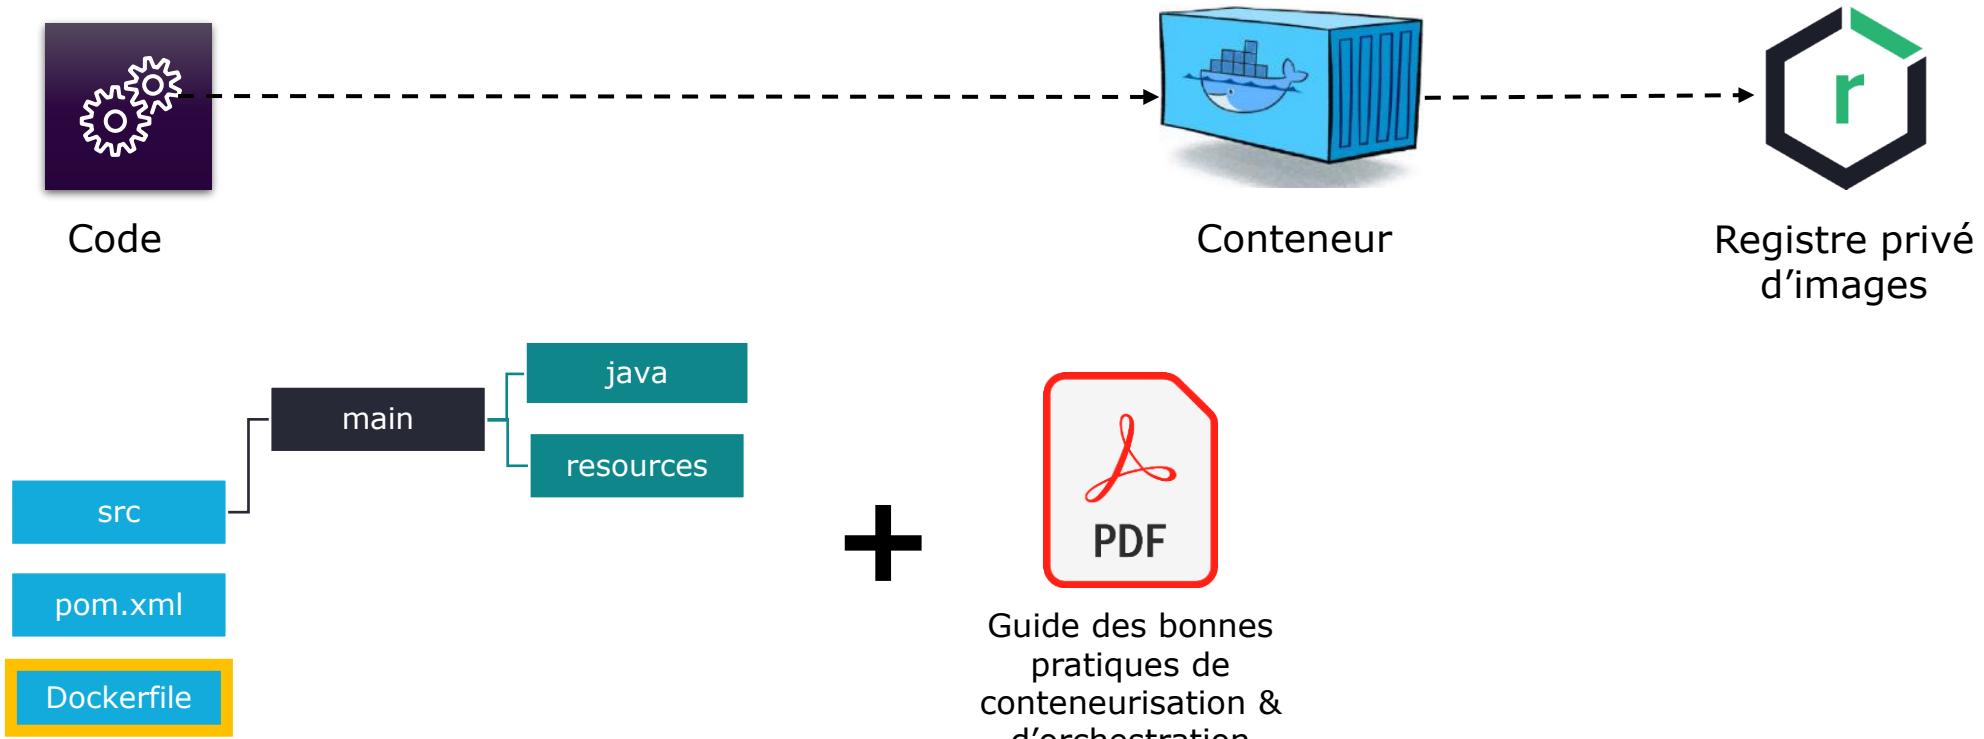
- Du code au cluster
- Automatiser la construction et le déploiement

Exploiter & superviser ses applications en production

- Exploiter ses applications
- Superviser ses applications
- Collecter des métriques

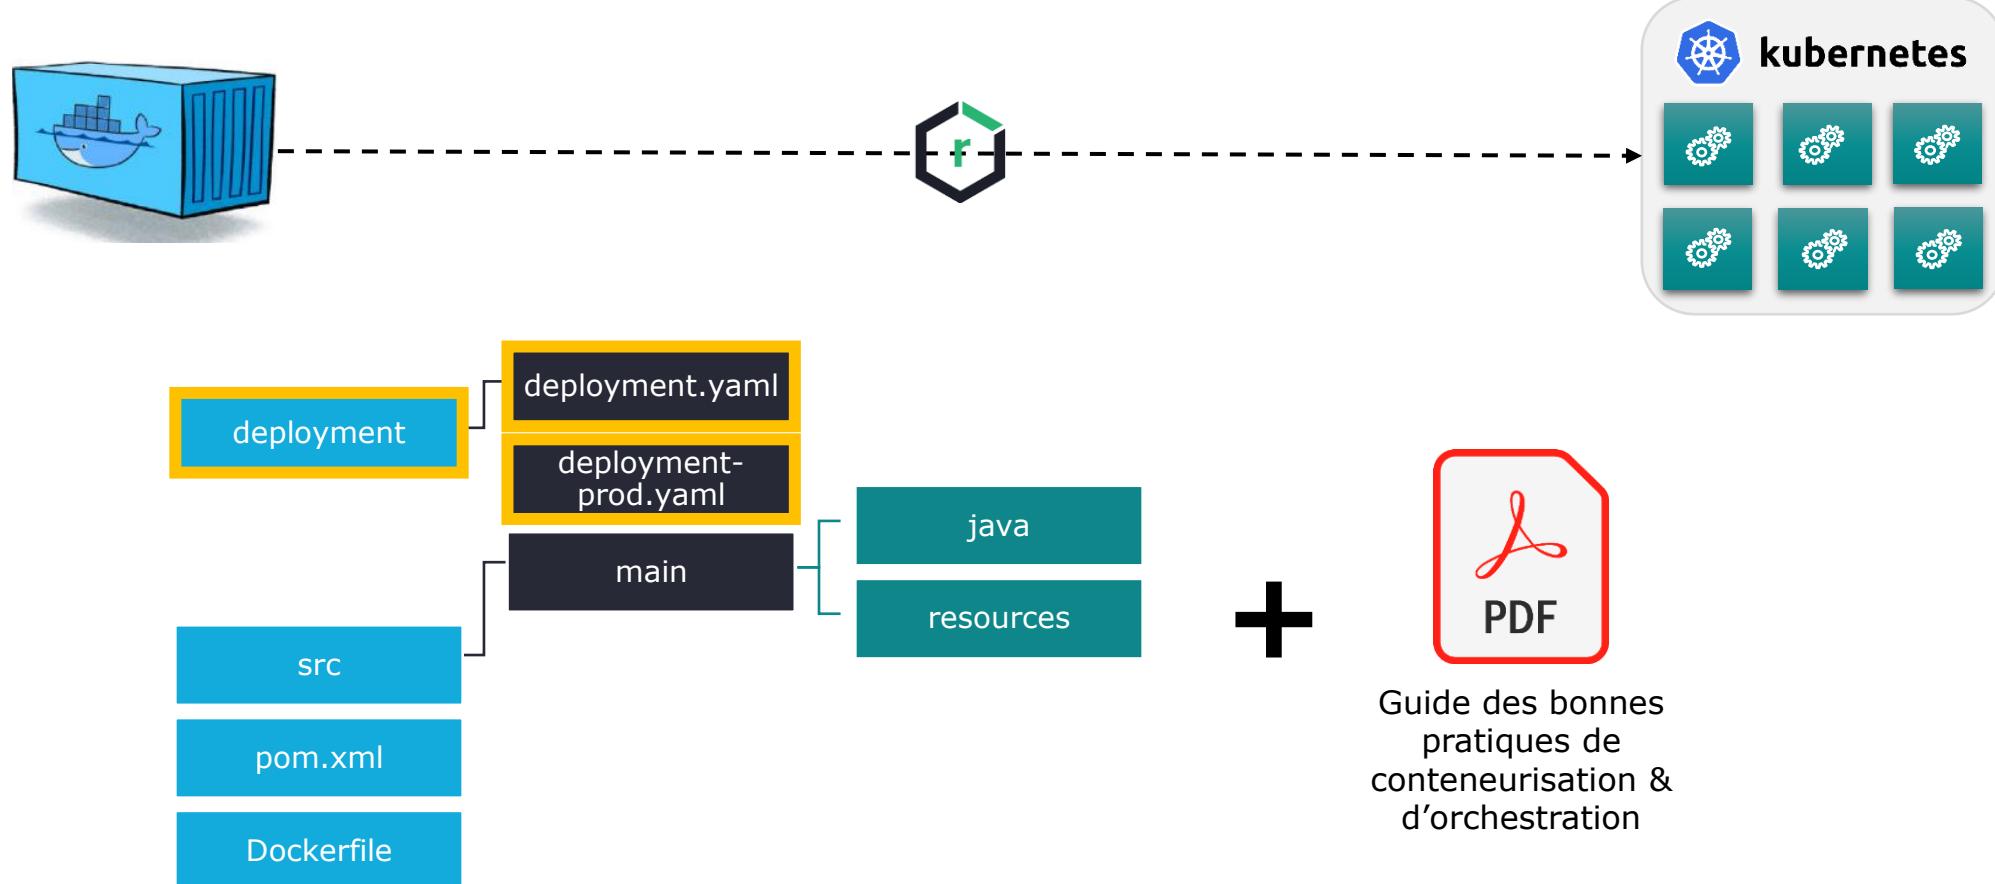
DU CODE AU CLUSTER

Packager son application au sein d'un conteneur



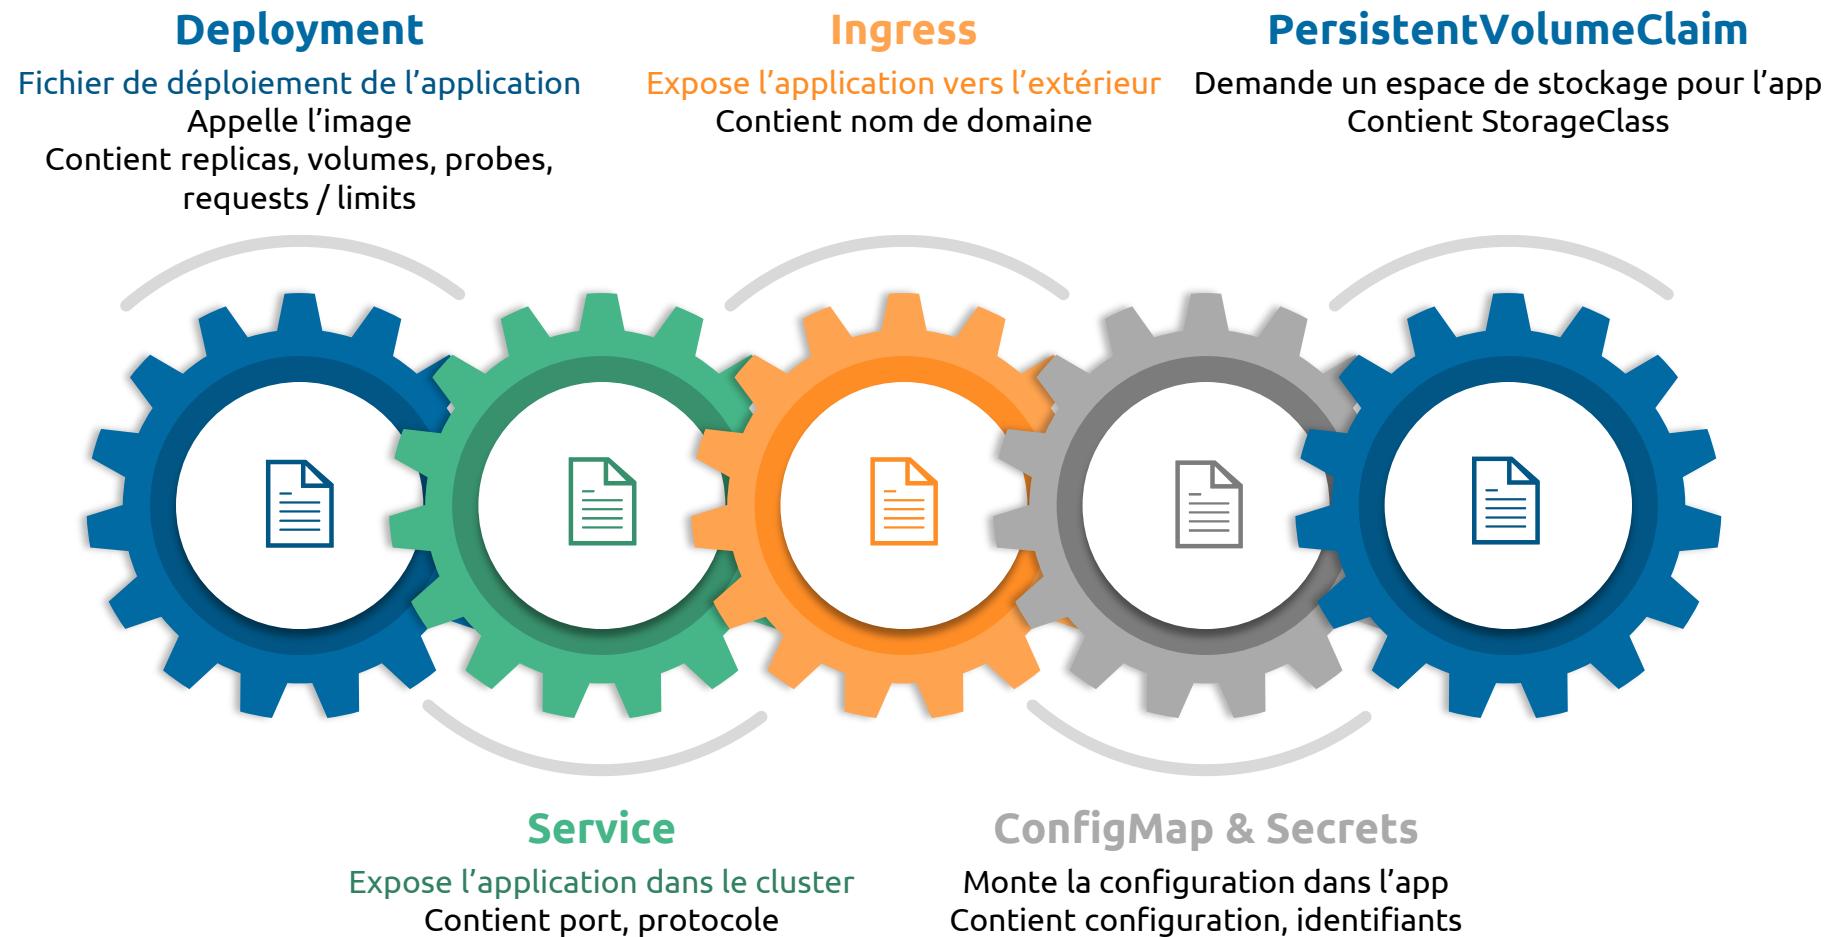
DU CODE AU CLUSTER

Héberger son conteneur sur des machines

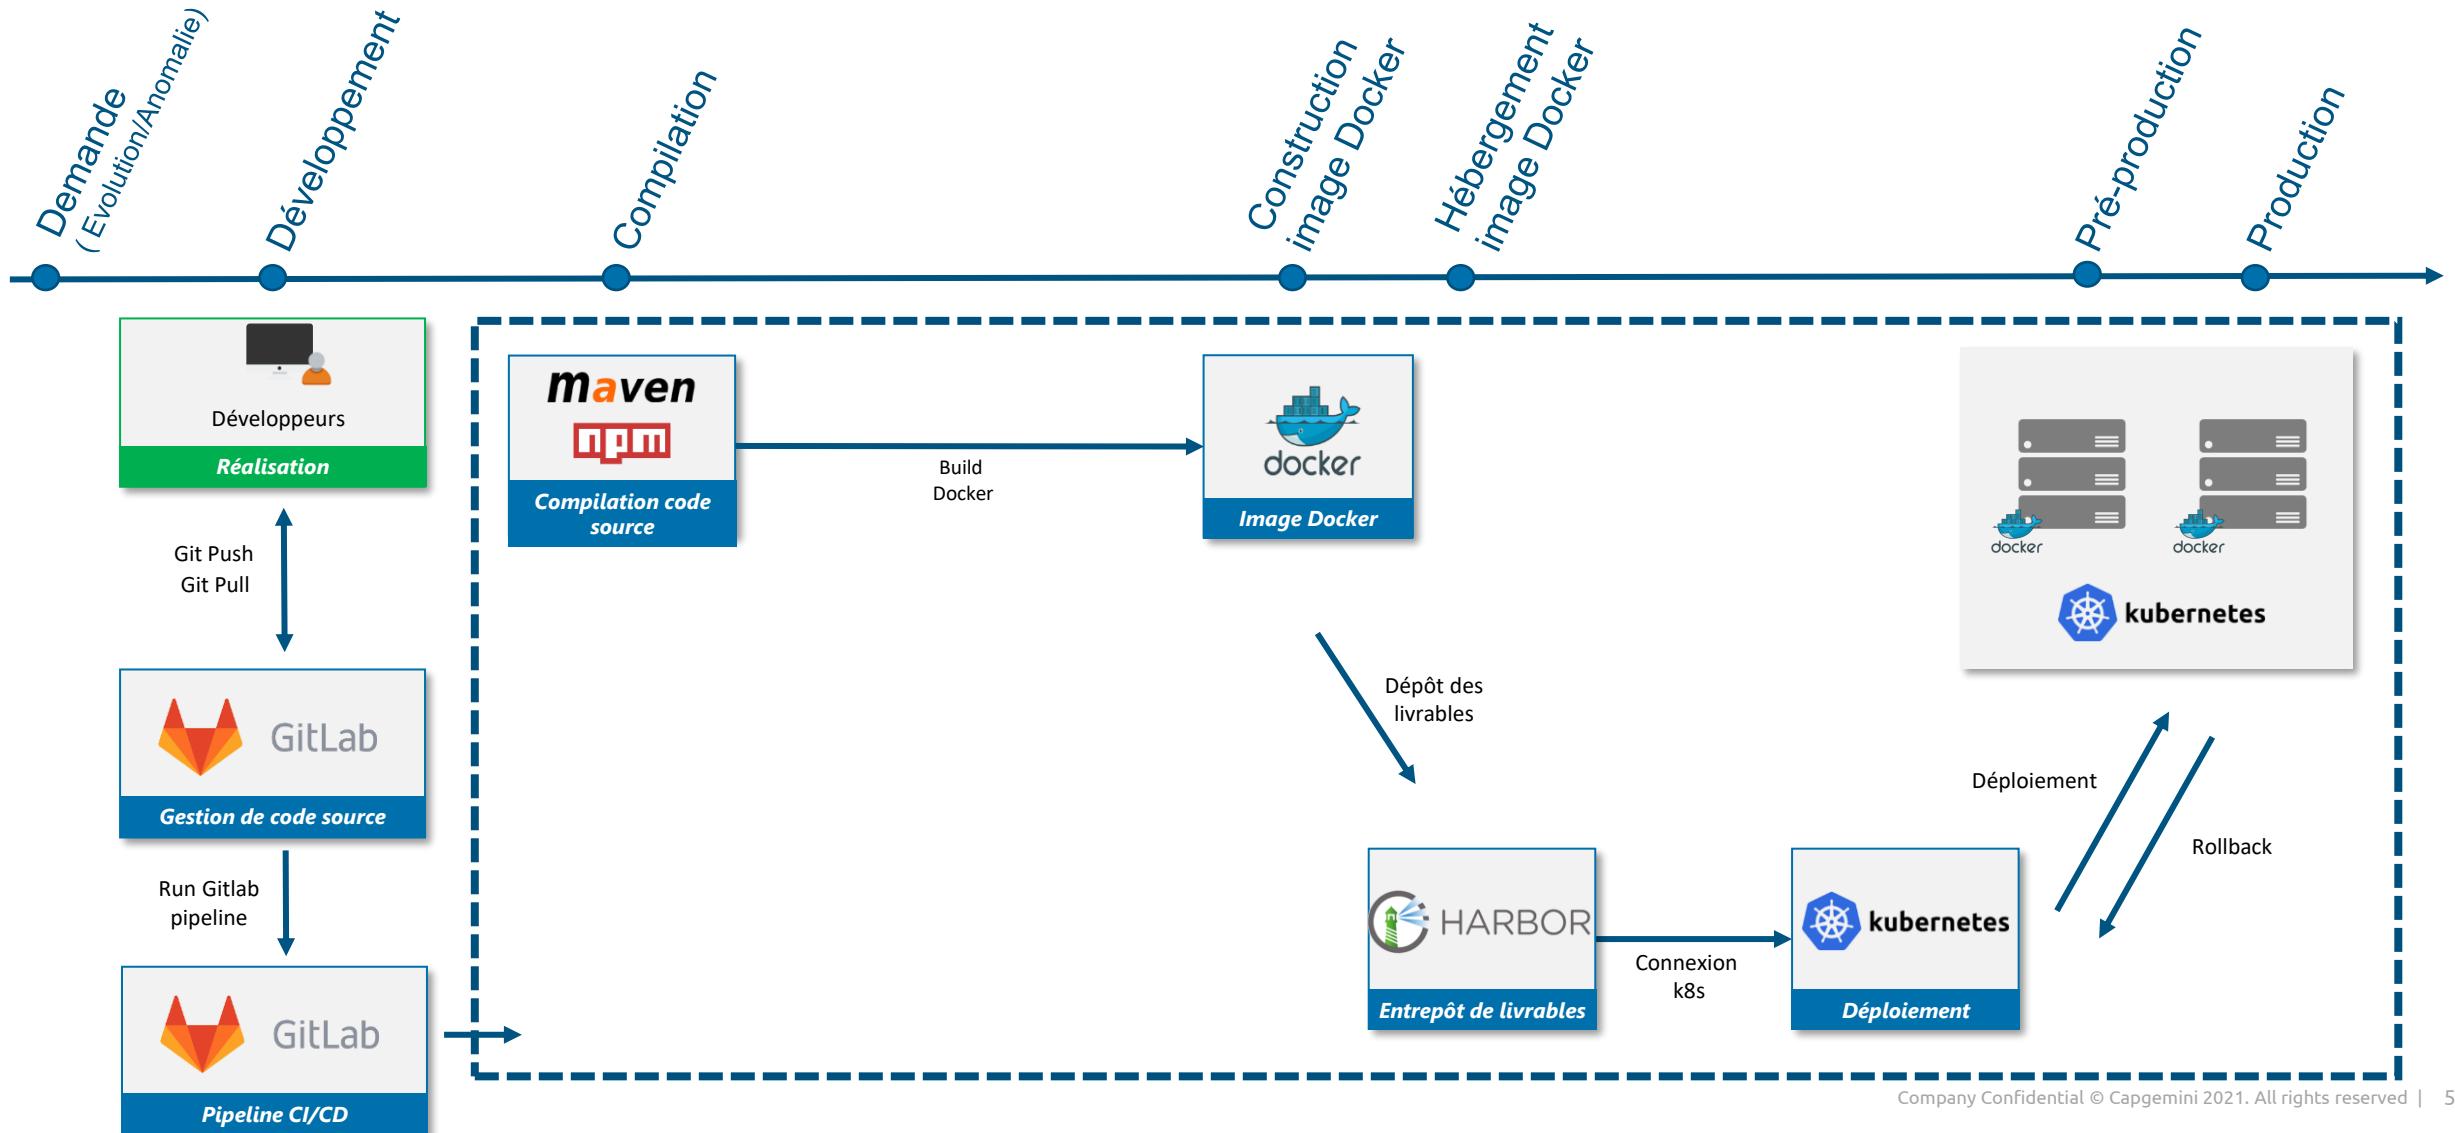


DU CODE AU CLUSTER

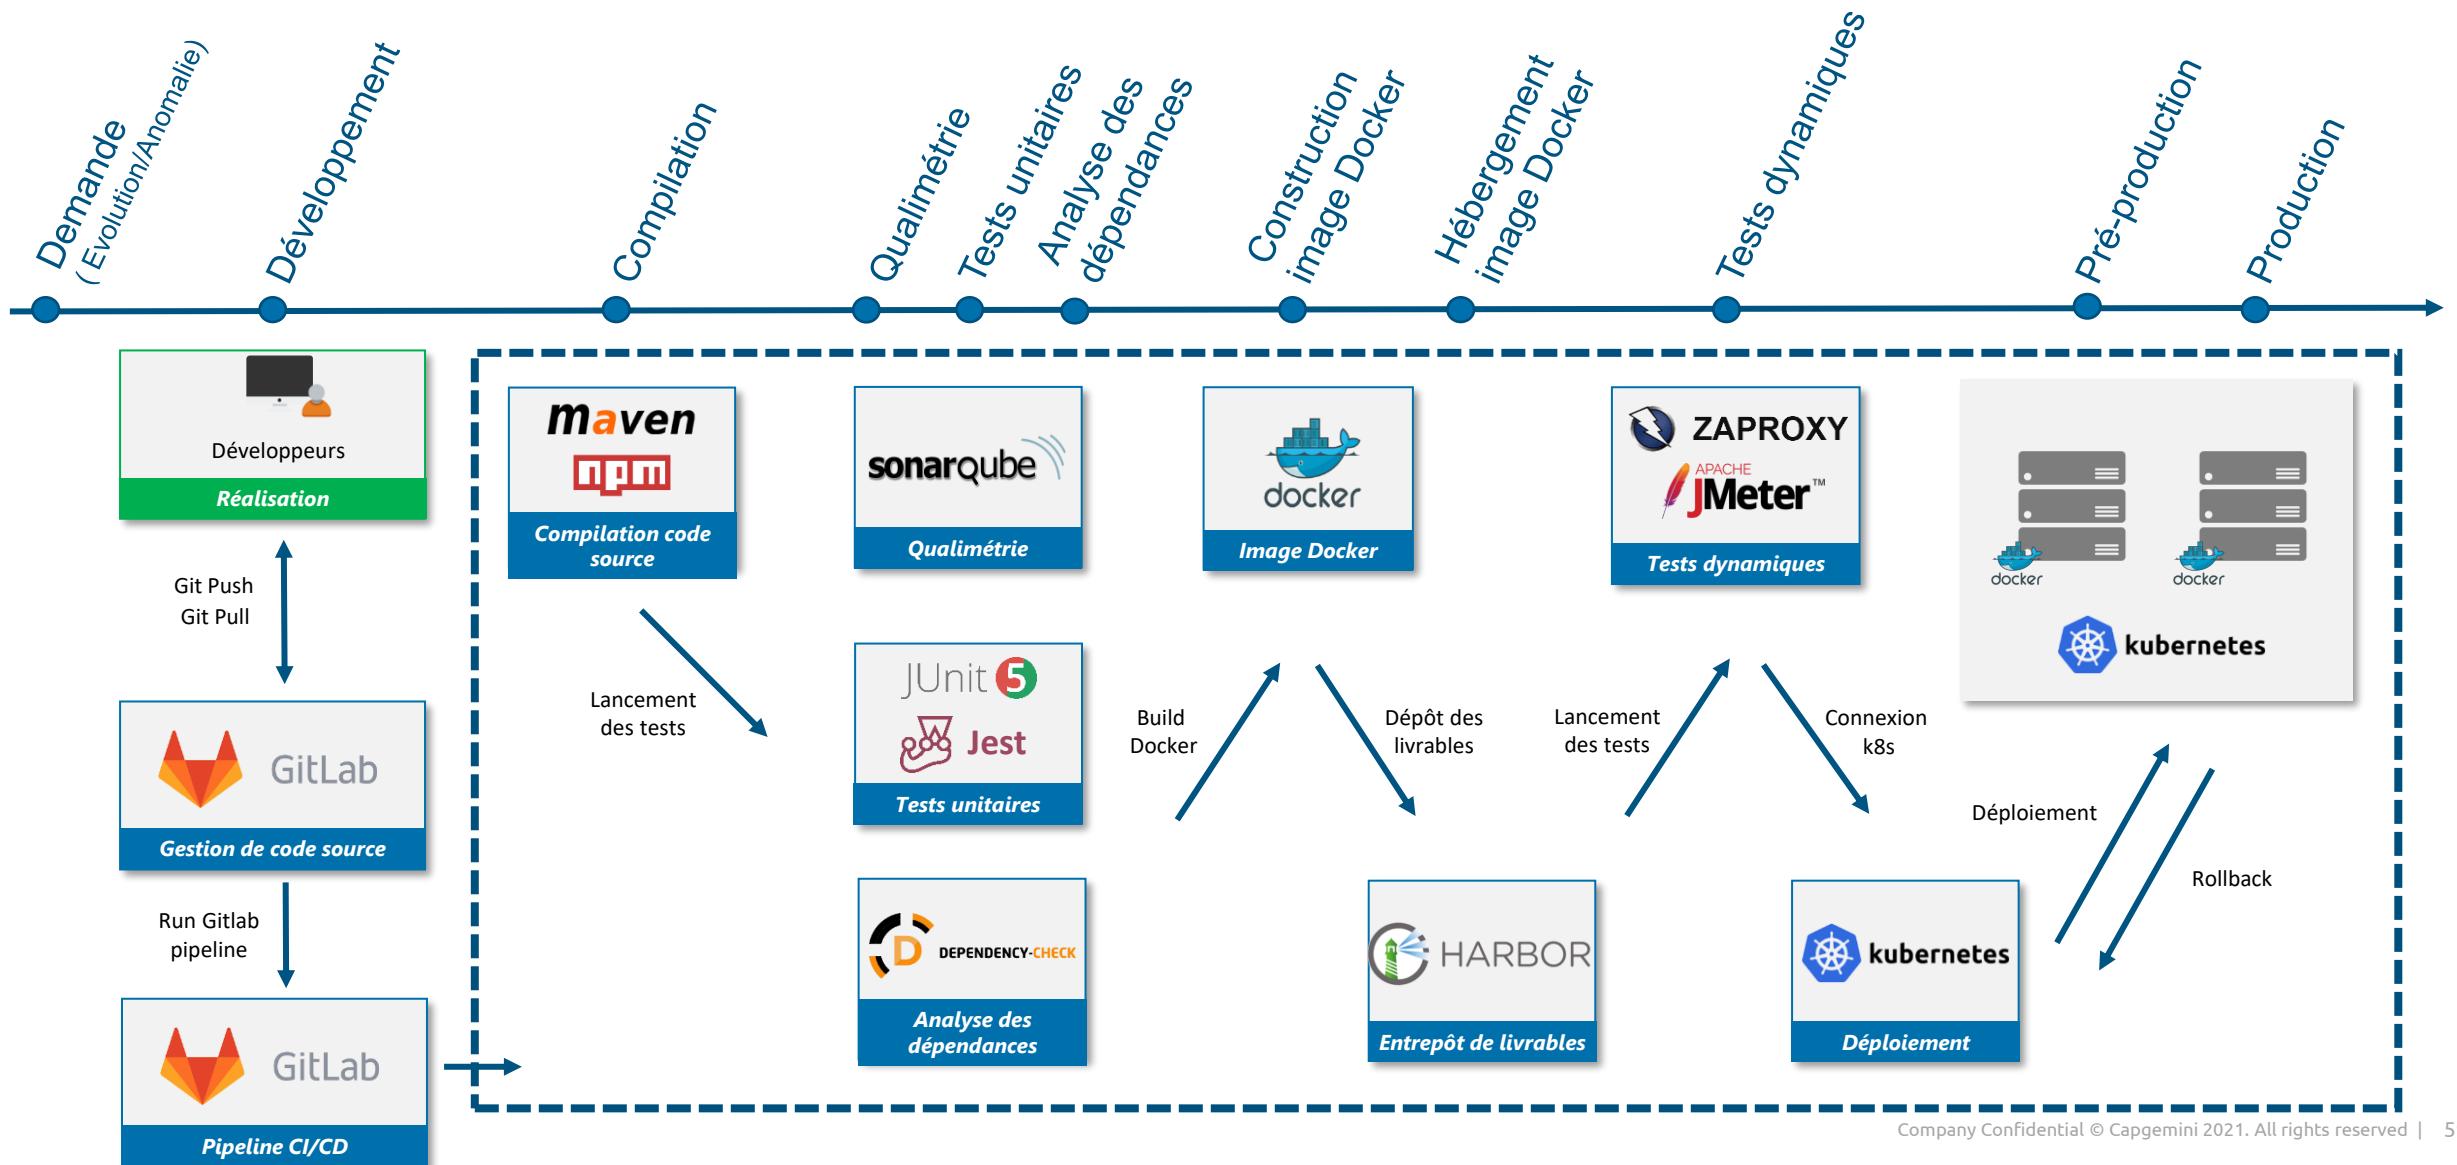
Adapter son application pour l'orchestrer dans Kubernetes



AUTOMATISER LA CONSTRUCTION ET LE DÉPLOIEMENT

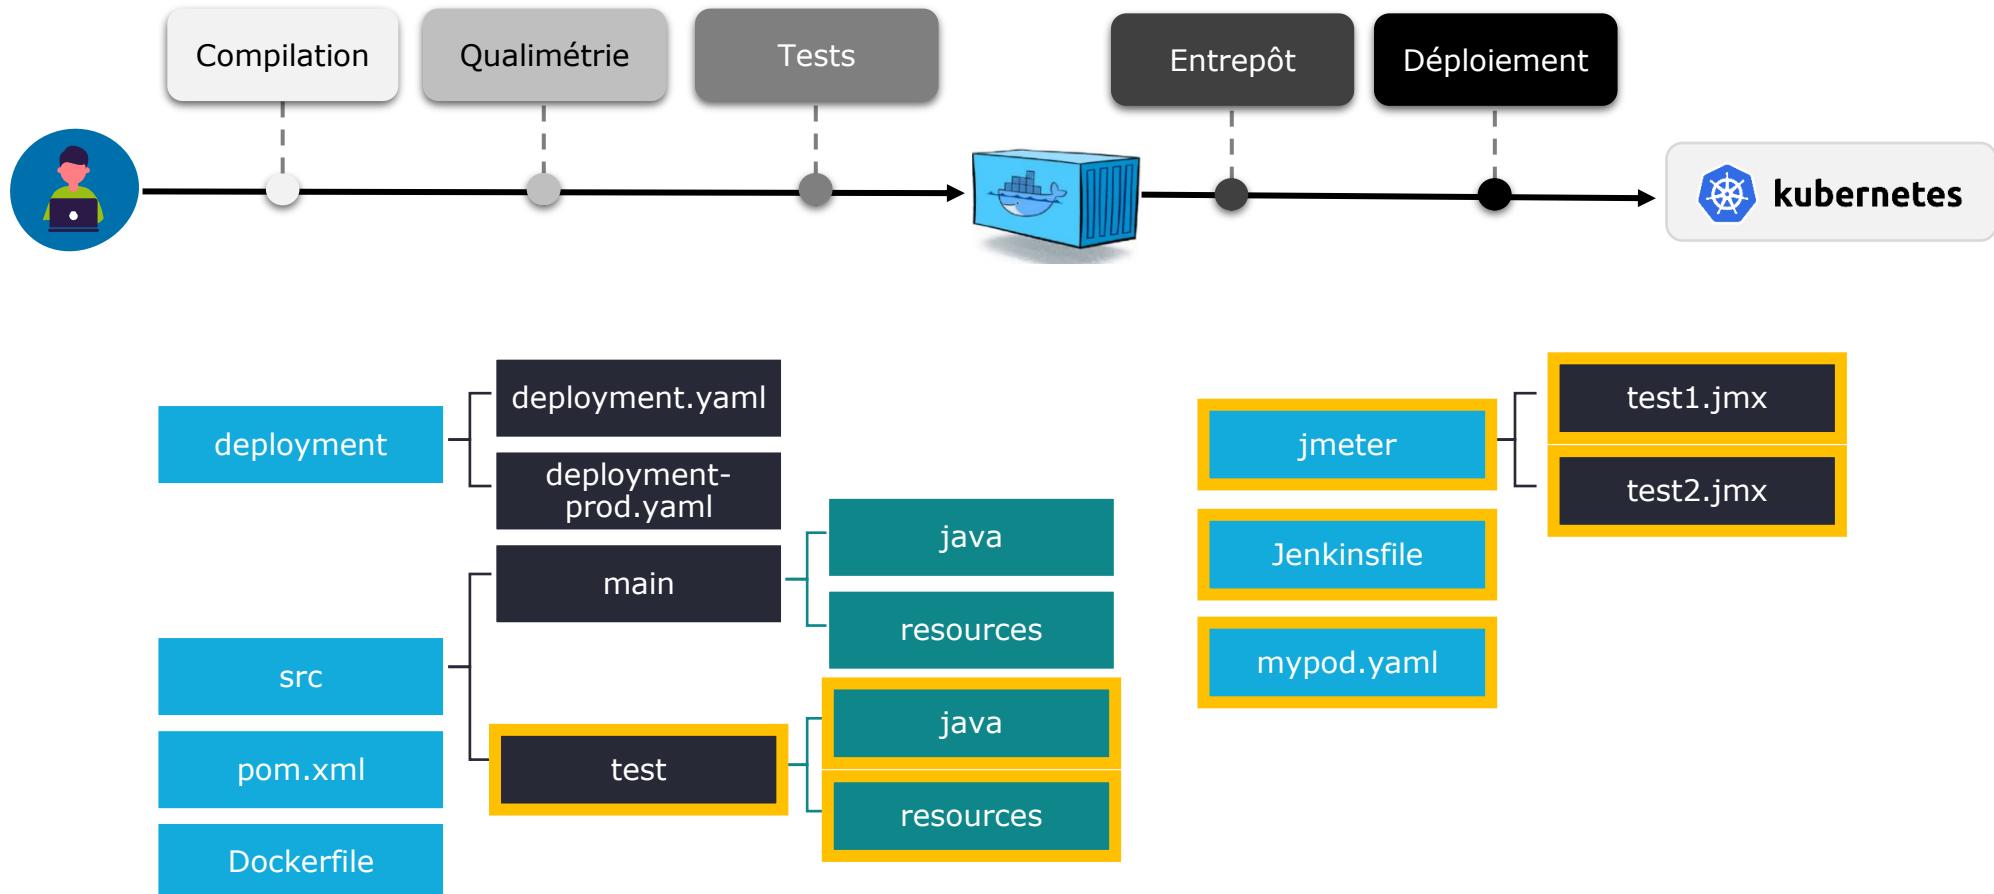


AUTOMATISER LA CONSTRUCTION ANALYTIQUE ET LE DÉPLOIEMENT



AUTOMATISER LA CONSTRUCTION ANALYTIQUE ET LE DÉPLOIEMENT

Adapter le contenu du dépôt de code





?

?

?



EXPLOITER SES APPLICATIONS EN PRODUCTION

Conteneuriser ses applications

- Concepts & bénéfices
- Création d'une image
- Lancement d'un conteneur

Orchestrer ses conteneurs via Kubernetes

- Concepts & bénéfices
- Déployer un cluster Kubernetes
- Les ressources Kubernetes
- Déployer des ressources sur le cluster

Industrialiser ses déploiements

- Du code au cluster
- Automatiser la construction et le déploiement

Exploiter & superviser ses applications en production

- Exploiter ses applications
- Superviser ses applications
- Collecter des métriques

EXPLOITER SES APPLICATIONS

Gérer le cycle de vie des applications en production

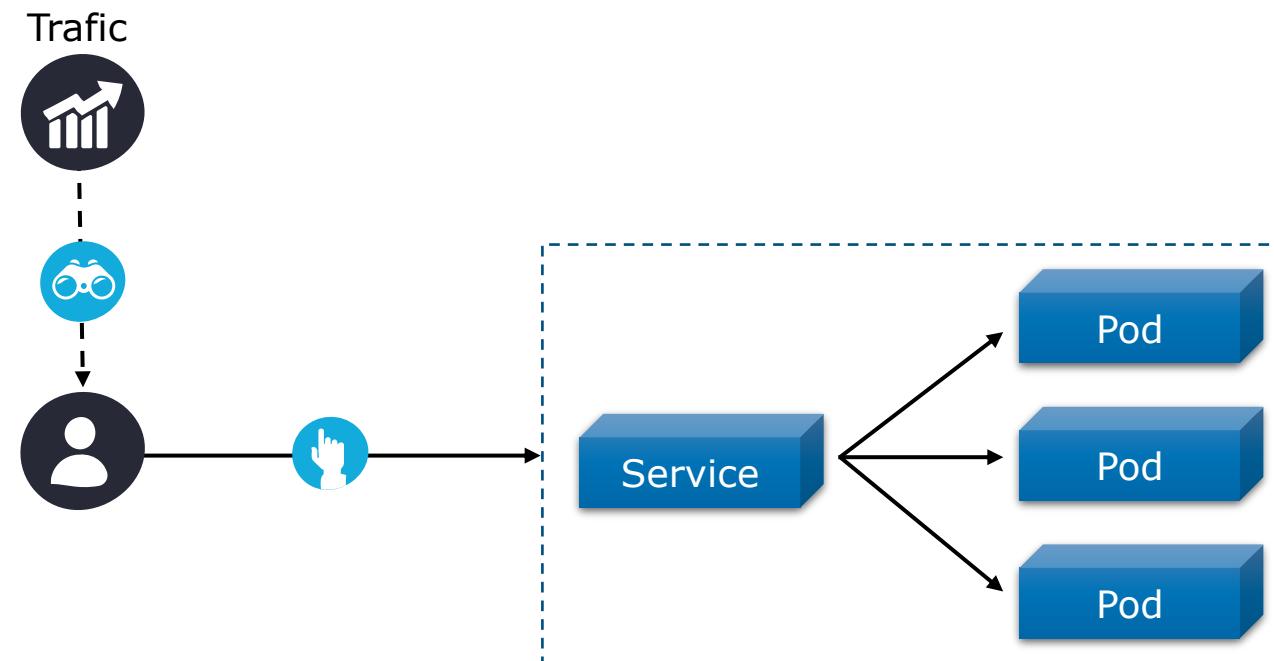
✓ Résilience

- Redémarrage en cas d'échec
- MAJ de la configuration
- Suppression de l'instance



✓ Passage à l'échelle

- Upscaling
- Downscaling



SUPERVISER SES APPLICATIONS

Intérêts de la supervision



Visualiser l'état du système

- Vérifier le fonctionnement
- Suivre les exigences
- Déetecter les incidents



Diagnostiquer & Investiguer

- Debugger
- Analyser les incidents (post-mortem)
- Analyser les tendances

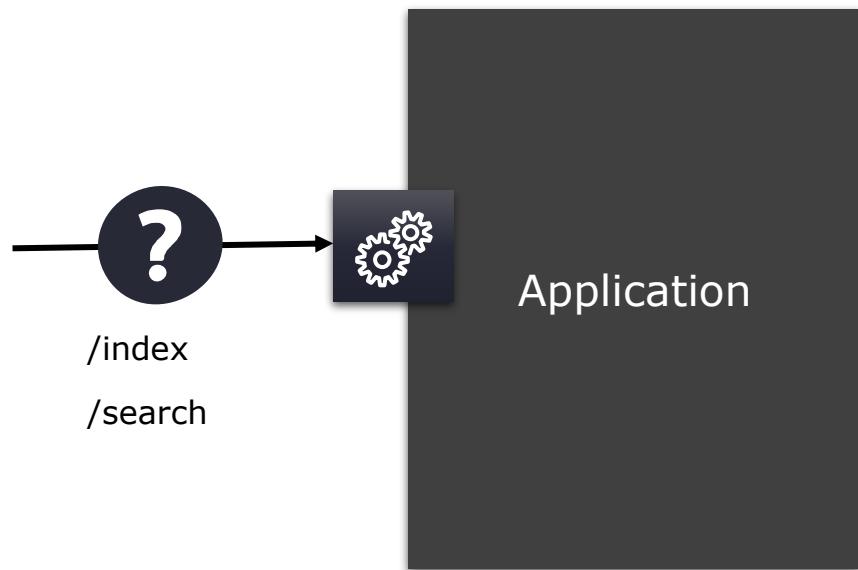


Corriger le système

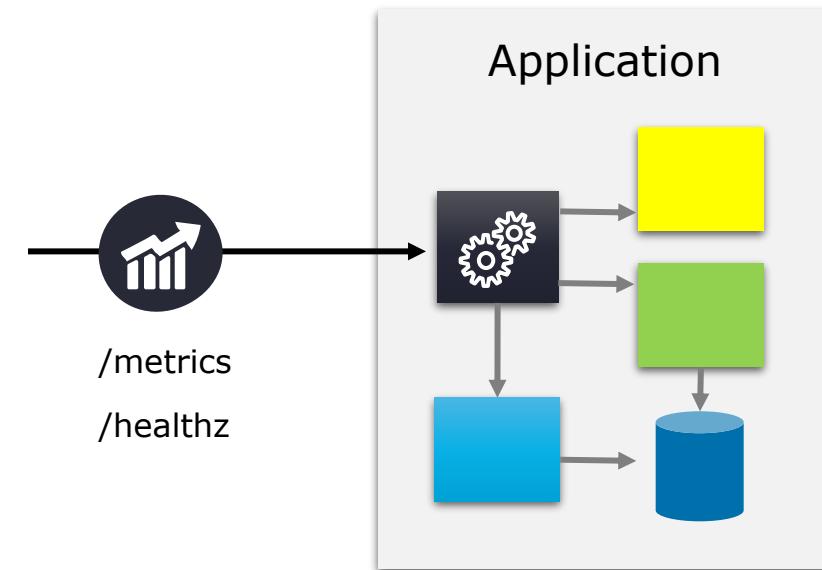
- Implémenter des correctifs
- Anticiper les pannes

SUPERVISER SES APPLICATIONS

Black-box vs white-box monitoring



Le service fonctionne-t-il ?



Le nombre important de requêtes risque de saturer le service.

SUPERVISER SES APPLICATIONS

Supervision vs observabilité

Observability is a measure of how well **internal states of a system** can be inferred from knowledge of its **external outputs**.

Wikipedia, <https://en.wikipedia.org/wiki/Observability>

Monitoring tells you **whether the system works**. Observability lets you ask **why it's not working**.

Baron Schwarz

OBSERVABILITÉ EN PRODUCTION

Métriques, logging et tracing : la supervision n'est plus une option !



Healthchecks



Logs



Metrics



Distributed Tracing



kubernetes



elasticsearch



Prometheus



JAEGER



Prometheus



kibana



Grafana loki



Grafana



kiali



OPENTRACING

https://blog.twitter.com/engineering/en_us/a/2013/observability-at-twitter.html

COLLECTER DES MÉTRIQUES

Métrique versus log

Métriques

Mesure numérique d'un fait pour un **instant donné**.

```
http_requests_total{code="200",handler="/",method="get"} 1
```

Logs

timestamp	2019-11-18T16:40:23+00:00		level	ERROR	log	Read time out
service	orders	team	alpha			
commit	548a8c9a	build	548a8c9a.3	runtime	java-1.8.0_164	
region	Paris-3	node	node_ea522aF28			
userId	389	requestId	ef524aC42	method	GET	endpoint /orders

COLLECTER DES MÉTRIQUES

Instrumenter ses développements



Coder des métriques customs



Instrumenter à l'aide d'une librairie

<https://prometheus.io/docs/instrumenting/clientlibs>



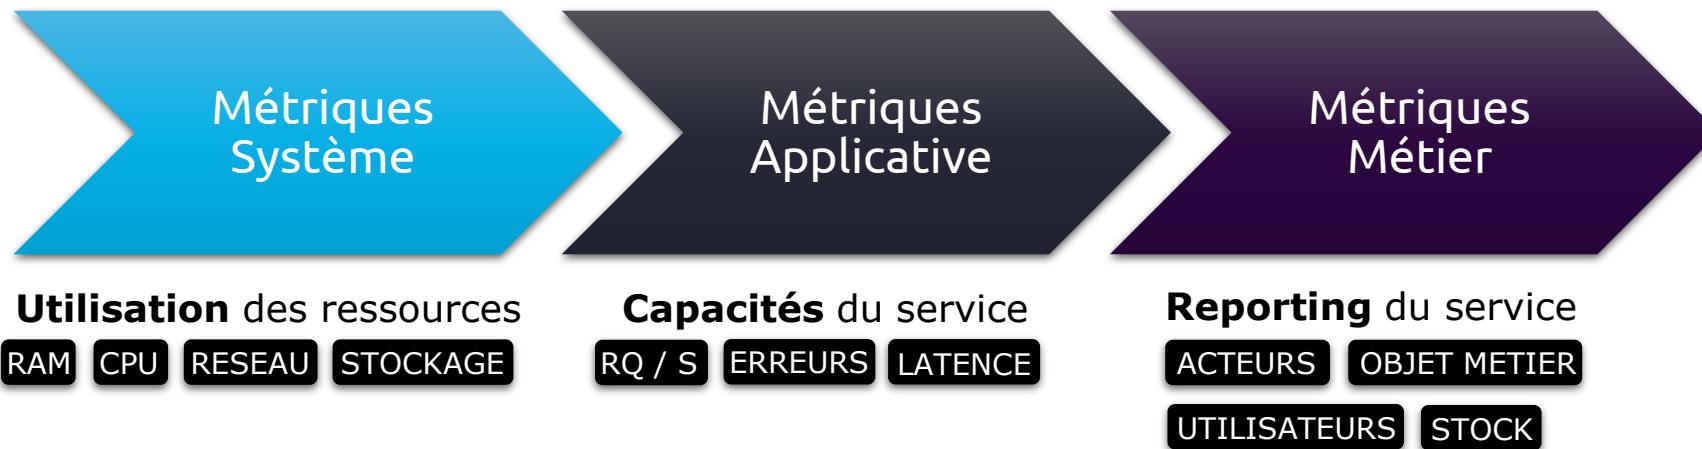
Utiliser un exporter / module existant

<https://prometheus.io/docs/instrumenting/exporters>



COLLECTER DES MÉTRIQUES

Typologie de métriques



DASHBOARD GRAFANA



DASHBOARD KIBANA

Dashboard / [Metricbeat Docker] Overview ECS

Full screen Share Clone Edit

Filters Search KQL Last 15 minutes Show dates Refresh

Docker containers [Metricbeat Docker] ECS

Name	CPU usage (%)	DiskIO	Mem (%)	Mem RSS	Number of Containers
docker_grafana_1	0.056	20.359	0.012	8.031MB	1
docker_apm_1	0.247	18.393	0.009	7.914MB	1
docker_db_1	0.528	12.194	0.009	4.586MB	1
docker_filebeat_1	0.856	9.573	0.008	6.613MB	1
docker_elasticsearch_1	17.161	1,034.124	0.358	673.805MB	1

18.848 1,094.643 0.396 5

Export: Raw Formatted

Number of Containers [Metricbeat Docker] ECS

9 0 20

Running Paused Stopped

Docker containers per host [Metricbeat Doc...]

Docker images and names [Metricbeat Docker] ECS

CPU usage [Metricbeat Docker] ECS

Count

18:48:00 18:49:00 18:50:00 18:51:00 18:52:00 18:53:00 18:54:00 18:55:00 18:56:00 18:57:00 18:58:00 18:59:00 19:00:00 19:01:00 19:02:00

@timestamp per 30 seconds

Memory usage [Metricbeat Docker] ECS

Count

18:48:00 18:49:00 18:50:00 18:51:00 18:52:00 18:53:00 18:54:00 18:55:00 18:56:00 18:57:00 18:58:00 18:59:00 19:00:00 19:01:00 19:02:00

@timestamp per 30 seconds

Network IO [Metricbeat Docker] ECS

Count

18:48:00 18:49:00 18:50:00 18:51:00 18:52:00 18:53:00 18:54:00 18:55:00 18:56:00 18:57:00 18:58:00 18:59:00 19:00:00 19:01:00 19:02:00

@timestamp per 30 seconds

Legend:

- docker_kibana_1: IN ... ● docker_kibana_1: OUT ... ● docker_elasticsearch_1: IN ... ● docker_elasticsearch_1: OUT ... ● docker_metricbeat_1: IN ... ● docker_metricbeat_1: OUT ... ● docker_apm_1: IN bytes ... ● docker_apm_1: OUT bytes ... ● docker_db_1: IN bytes ... ● docker_db_1: OUT bytes ... ● docker_grafana_1: IN ... ● docker_grafana_1: OUT ...
- docker_web_1: IN bytes ... ● docker_web_1: OUT bytes ... ● docker_logstash_1: IN ... ● docker_logstash_1: OUT ...

292.969KB

TRACES KIBANA

Trace sample < 1 of 500 >

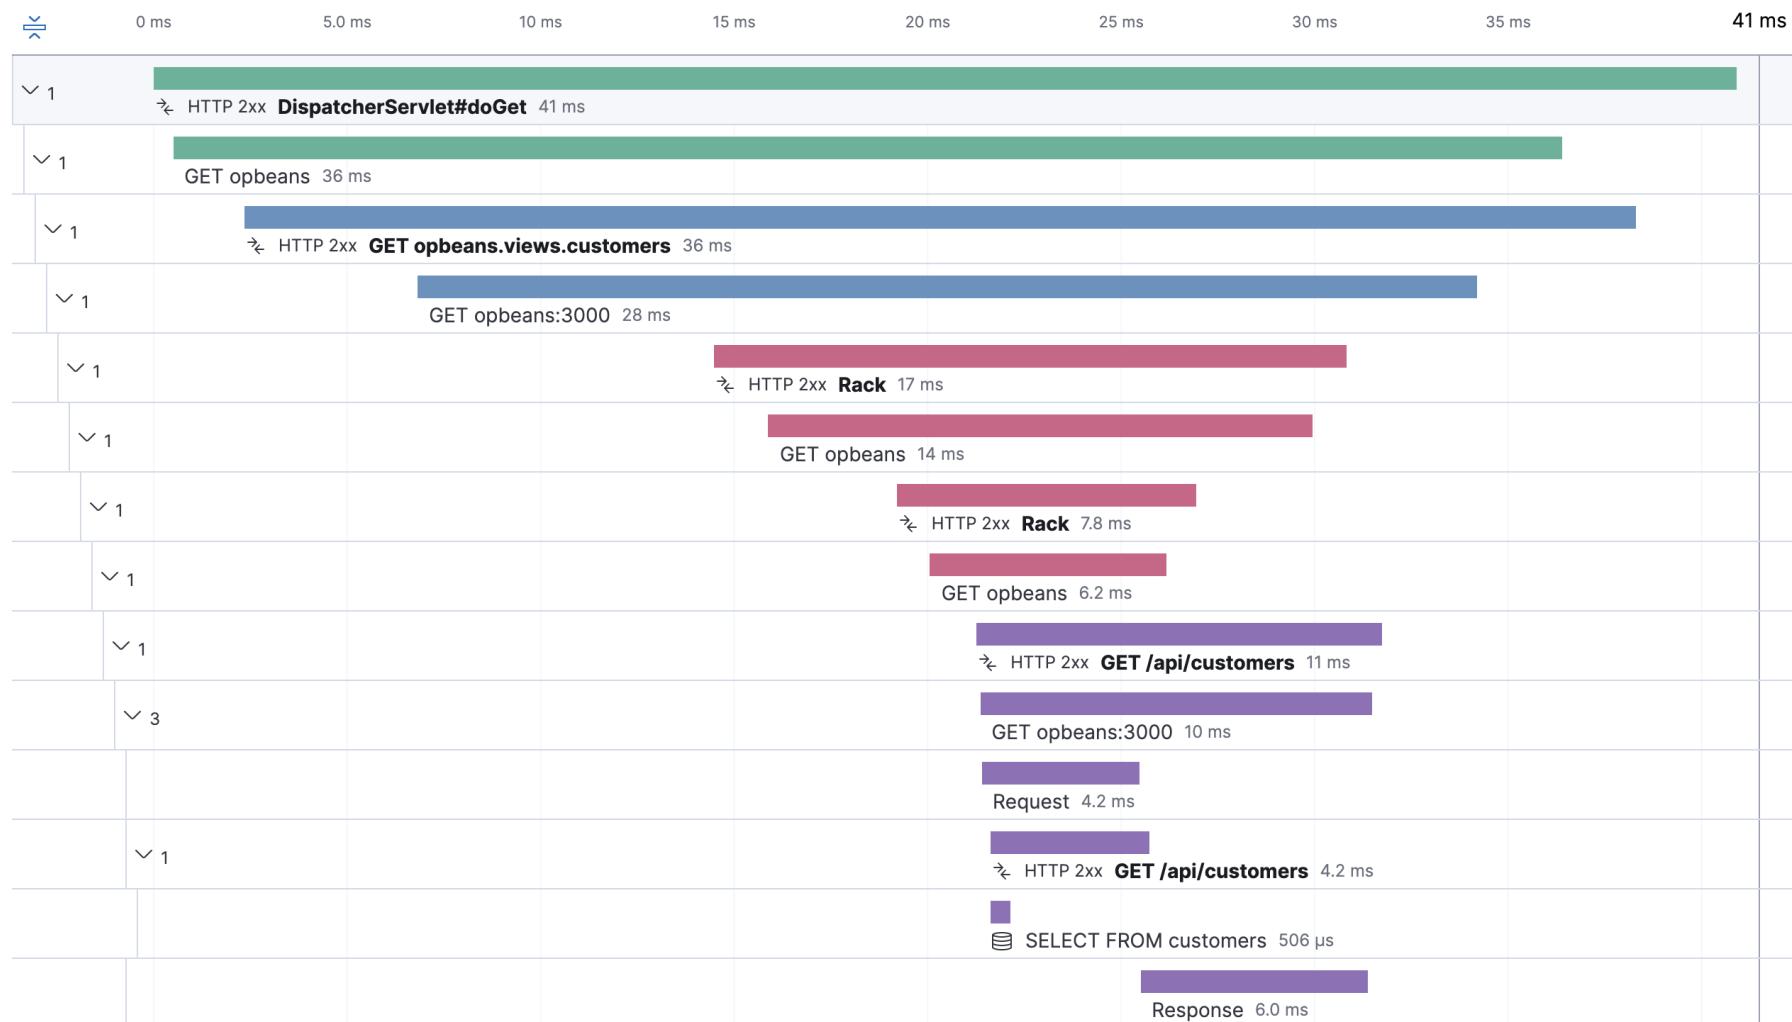
Investigate ▾

View full trace

2 days ago | 41 ms (98% of trace) | GET http://10.11.249.198:3000/api/customers | 200 OK | Python aiohttp (3.3.2)

Timeline Metadata Logs

Services ● opbeans-java ● opbeans-python ● opbeans-ruby ● opbeans-go



MONITORING KIBANA

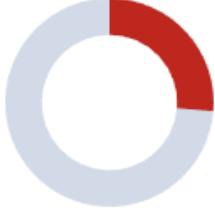
≡ |  |  | Uptime |  |  | 

Overview

Alerts ▾  Settings  Last 15 minutes Show dates Refresh

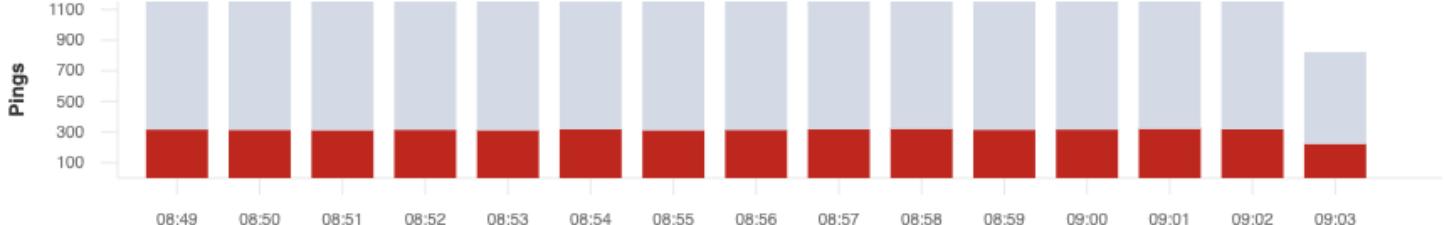
Search monitor IDs, names, and protocol types... Location 3 Port 7 Scheme 3 Tags 1

38 Monitors



● Down	10
● Up	28

Pings over time



Time	Successful Pings	Total Pings
08:49	~300	~1100
08:50	~300	~1100
08:51	~300	~1100
08:52	~300	~1100
08:53	~300	~1100
08:54	~300	~1100
08:55	~300	~1100
08:56	~300	~1100
08:57	~300	~1100
08:58	~300	~1100
08:59	~300	~1100
09:00	~300	~1100
09:01	~300	~1100
09:02	~300	~1100
09:03	~300	~800

Monitors

All Up Down Certificates status

Status	Name	Url	TLS Certificate	Downtime history
● Down 5s ago	in 3/3 Locations	Website Monitor - Infra Error	https://www.elastic.co/products/in...  Expires in 2 months	
● Up 4s ago	in 3/3 Locations	NodeJS	http://opbeans-node:3000/api/cus... -	
● Up 4s ago	in 3/3 Locations	NodeJS	http://opbeans-node:3000/api/sta... -	
● Up 8s ago	in 3/3 Locations	NodeJS	http://opbeans-node:3000/api/ord... -	



?

?

?

RECOMMANDATIONS

Comprendre le Kube

Kubernetes for everyone



#kubernetes #architecture #beginners #docker

 Sendil Kumar N

16 juil. 2019 Originally published at sendilkumarn.com · Updated on 27 nov. 2020 · 7 min read

<https://dev.to/sendilkumarn/kubernetes-for-everyone-opb>

Understanding kubernetes networking: pods

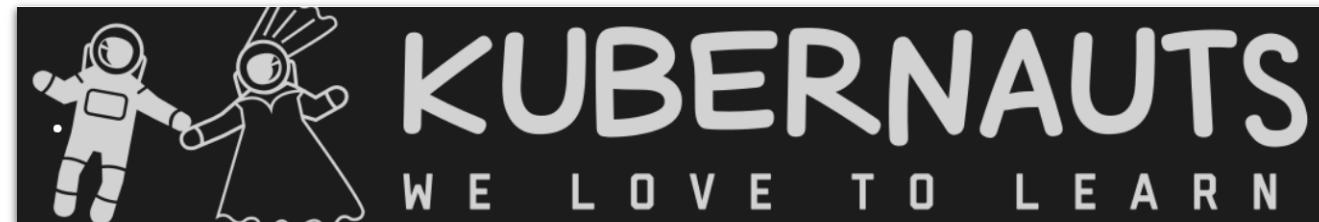


Mark Betz [Follow](#)

Oct 27, 2017 · 9 min read



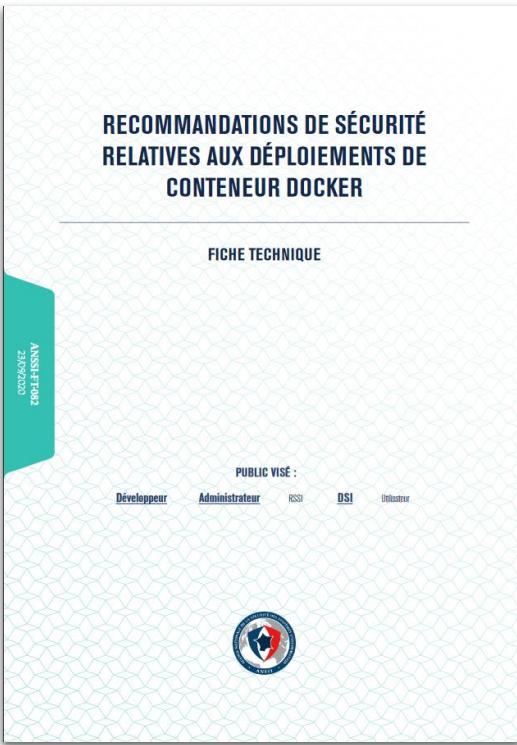
<https://medium.com/google-cloud/understanding-kubernetes-networking-pods-7117dd28727>



<https://blog.kubernauts.io/>

RECOMMANDATIONS

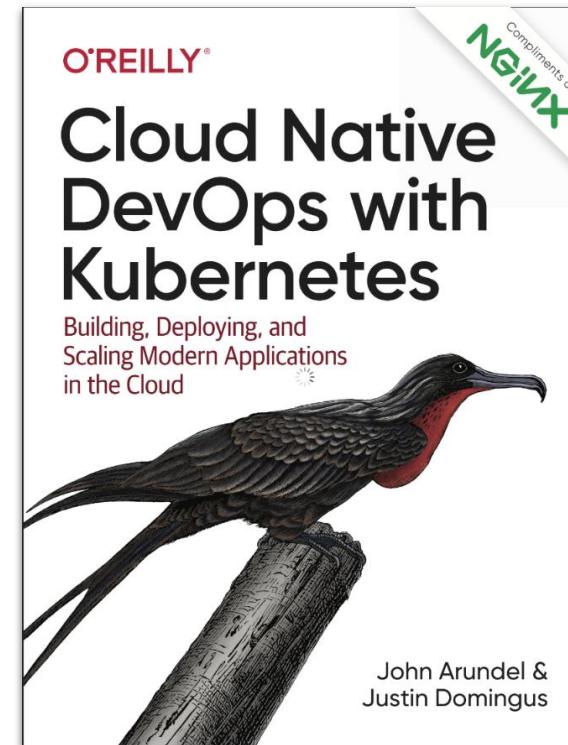
Implémenter le Kube



<https://www.ssi.gouv.fr/guide/recommandations-de-securite-relatives-aux-deploiement-de-conteneurs-docker/>



<https://learnk8s.io/production-best-practices>



<https://www.oreilly.com/library/view/cloud-native-devops/9781492040750/>



MERCI POUR VOTRE ATTENTION



About Capgemini

Capgemini is a global leader in partnering with companies to transform and manage their business by harnessing the power of technology. The Group is guided everyday by its purpose of unleashing human energy through technology for an inclusive and sustainable future. It is a responsible and diverse organization of 270,000 team members in nearly 50 countries. With its strong 50 year heritage and deep industry expertise, Capgemini is trusted by its clients to address the entire breadth of their business needs, from strategy and design to operations, fuelled by the fast evolving and innovative world of cloud, data, AI, connectivity, software, digital engineering and platforms. The Group reported in 2020 global revenues of €16 billion.



Get the Future You Want | www.capgemini.com

This presentation contains information that may be privileged or confidential and is the property of the Capgemini Group.

Copyright © 2021 Capgemini. All rights reserved.